

ISE 5103 Intelligent Data Analytics

Homework 5 - Modeling

Daniel Carpenter & Sonaxy Mohanty

October 2022

Contents

Packages	2
General Data Prep	2
Read Data	2
Impute Missing Values with PMM	2
Factor Level Collapse - Create Other Bin for Columns over 4 Unique Values	5
Remove Outliers from Numeric Data	6
1 (a) - OLS Model	7
1 (b) - PLS Model	8
1 (c) - LASSO Model	9
1 (d) - Model Variants	10
1 (d, i) - PCR Model	10
Perform PCA analysis to see how Principal components explain variance	10
Now, Apply predictions with PCR	11
Interpretation of PCR Model	13
Visualizatoion of PCR Model - Predicted vs. Actuals	15
1 (d, ii) - SVR Model	18
Model Setup	18
Tune Model	18
Diagnostics of Model	18
1 (d, iii) - MARS Model	21
Summary Table of Model Performance	22

Packages

```
# Data Wrangling
library(tidyverse)

# Modeling
library(MASS)
library(caret) # Modeling variants like SVM

# Aesthetics
library(knitr)
library(cowplot) # multiple ggplots on one plot with plot_grid()
library(scales)
library(kableExtra)
```

General Data Prep

Read Data

```
hd <- read.csv('housingData.csv') %>%

# creates new variables age, ageSinceRemodel, and ageofGarage, and
dplyr::mutate(age = YrSold - YearBuilt,
              ageSinceRemodel = YrSold - YearRemodAdd,
              ageofGarage = ifelse(is.na(GarageYrBlt), age, YrSold - GarageYrBlt)) %>%

# removes the columns used in above the calculations
dplyr::select(!c(Id, MSSubClass, MiscVal, YrSold,
                 MoSold, YearBuilt, YearRemodAdd))

# Convert all character data to factor
hd[sapply(hd, is.character)] <-
  lapply(hd[sapply(hd, is.character)], as.factor)
```

Impute Missing Values with PMM

Make dataset of **numeric** variables

```
hd.numericRaw <- hd %>%

#selecting all the numeric data
dplyr::select_if(is.numeric) %>%

#converting the dataframe to tibble
as_tibble()
```

Make dataset of **factor** variables

```

hd.factorRaw <- hd %>%

#selecting all the numeric data
dplyr::select_if(is.numeric) %>%

#converting the dataframe to tibble
as_tibble()

```

For each column with missing data, impute missing values with PMM

- Done with function `imputeWithPMM()` function
- Applies function via `dplyr` logic
- Note `seeImputation()` function to visualize the imputation from prior homework 4, not shown for simplicity in viewing

Create function to impute via PMM

```

imputeWithPMM <- function(colWithMissingData) {

# Using the mice package
suppressMessages(library(mice))

# Discover the missing rows
isMissing <- is.na(colWithMissingData)

# Create data frame to pass to PMM imputation function from mic package
df <- data.frame(x      = rexp(length(colWithMissingData)), # meaningless x to help show variation
                 y      = colWithMissingData,
                 missing = isMissing)

# imputation by PMM
df[isMissing, "y"] <- mice.impute.pmm( df$y,
                                       !df$missing,
                                       df$x)

return(df$y)
}

```

Apply PMM function to numeric data containing null values

```

# Data to store imputed values with PMM method
hd.Imputed <- hd

# Which columns has NA's?
colNamesWithNulls <- colnames(hd.numericRaw[ , colSums(is.na(hd.numericRaw)) != 0])
colNamesWithNulls

```

```
## [1] "LotFrontage" "MasVnrArea" "GarageYrBlt"
```

```

numberOfColsWithNulls = length(colNamesWithNulls)

# For each of the numeric columns with null values
for (colWithNullsNum in 1:numberOfColsWithNulls) {

  # The name of the column with null values
  nameOfThisColumn <- colNamesWithNulls[colWithNullsNum]

  # Get the actual data of the column with nulls
  colWithNulls <- hd[, nameOfThisColumn]

  # Impute the missing values with PMM
  imputedValues <- imputeWithPMM(colWithNulls)

  # Now store the data in the original new frame
  hd.Imputed[, nameOfThisColumn] <- imputedValues

  # Save a visualization of the imputation
  pmmVisual <- seeImputation(data.frame(y = colWithNulls),
                             data.frame(y = imputedValues),
                             nameOfThisColumn )

  fileToSave = paste0('OutputPMM/Imputation_With_PMM_', nameOfThisColumn, '.pdf')
  print(paste0('For imputation results of ', nameOfThisColumn, ', see ', fileToSave))
  ggsave(pmmVisual, filename = fileToSave,
         height = 11, width = 8.5)
}

```

```
## [1] "For imputation results of LotFrontage, see OutputPMM/Imputation_With_PMM_LotFrontage.pdf"
```

```
## [1] "For imputation results of MasVnrArea, see OutputPMM/Imputation_With_PMM_MasVnrArea.pdf"
```

```
## [1] "For imputation results of GarageYrBlt, see OutputPMM/Imputation_With_PMM_GarageYrBlt.pdf"
```

Factor Level Collapse - Create Other Bin for Columns over 4 Unique Values

```
hd.Cleaned <- hd.Imputed # For final cleaned data

# Get list of factors and the number of unique values
factorCols <- as.data.frame(t(hd.factorRaw %>% summarise_all(n_distinct)))

# We are going to factor collapse factor columns with more than 4 columns
# So there will be 4 of the original, and 1 containing 'other'
# This is the threshold
factorThreshold = 4

# Get a list of the factors we are going to collapse
colsWithManyFactors <- rownames(factorCols %>% filter(V1 > factorThreshold))

# Show a summary of how many factors will be collapsed
numberOfColsWithManyFactors = length(colsWithManyFactors)
paste('Before cleaning, there are', numberOfColsWithManyFactors, 'factor columns with more than',
      factorThreshold, 'unique values')
```

```
## [1] "Before cleaning, there are 14 factor columns with more than 4 unique values"
```

```
# Collapse the affected factors in the original data (the one that already has imputation)

## for each factor column that we are about to collapse
for (collapsedColNum in 1:numberOfColsWithManyFactors) {

  # The name of the column with null values
  nameOfThisColumn <- colsWithManyFactors[collapsedColNum]

  # Get the actual data of the column with nulls
  colWithManyFactors <- hd[, nameOfThisColumn]

  # lumps all levels except for the n most frequent
  hd.Cleaned[, nameOfThisColumn] <- fct_lump_n(colWithManyFactors,
                                              n=factorThreshold)
}

# Check to see if the factor lumping worked
factorColsCleaned <- t(hd.Cleaned %>%
  select_if(is.factor) %>%
  summarise_all(n_distinct))
paste('After cleaning, there are', sum(factorColsCleaned > factorThreshold, na.rm = TRUE),
      "columns with more than", factorThreshold, "unique values (omitting NA's)")
```

```
## [1] "After cleaning, there are 14 columns with more than 4 unique values (omitting NA's)"
```

Remove Outliers from Numeric Data

- Since there are so many outliers, we are only going to remove some outliers
- If you count the number of outliers by column, the 75% of columns contain less than 50 outliers.
- However, some contain up to 200. Since remove ALL outliers would reduce the size of the data to less than 300 observations, we are removing up to 50 per column.

```
hd.CleanedNoOutliers <- hd.Cleaned

# Remove up to 75% of the outliers in the dataset
# this is the 3rd quartile of number of outliers.
k_outliers = 50
numOutliers = c() # to store the number of outliers per column

theColNames <- colnames(hd.Cleaned)

for (colNum in 1:ncol(hd.Cleaned)) {

  theCol <- hd.Cleaned[, colNum]
  nrowBefore = length(theCol)
  colName <- theColNames[colNum]

  # Only consider numeric
  if (is.numeric(theCol)) {

    # Identify the outliers in the column
    # Source: https://www.geeksforgeeks.org/remove-outliers-from-data-set-in-r/
    columnOutliers <- boxplot.stats(hd.CleanedNoOutliers[, colNum])$out
    numOutliers <- c(numOutliers, length(columnOutliers))

    # Now remove k outliers from the column
    if (length(columnOutliers) < k_outliers) {

      hd.CleanedNoOutliers <- hd.CleanedNoOutliers %>%

        # If this syntax looks weird, it is just referencing a column in the
        # dataset using dplyr piping. See below for more info:
        # https://stackoverflow.com/questions/48062213/dplyr-using-column-names-as-function-arguments
        # https://stackoverflow.com/questions/72673381/column-names-as-variables-in-dplyr-select-v-filter
        filter( !( get({colName}) ) %in% columnOutliers ) )
    }
  }
}

paste0('Of the columns with outliers, removed up to 75th percentile of num. outliers.')

## [1] "Of the columns with outliers, removed up to 75th percentile of num. outliers."

paste0('See that the 75th percentile of columns with outliers contain ',
       paste0(summary(numOutliers)[5]), ' outliers')

## [1] "See that the 75th percentile of columns with outliers contain 51.25 outliers"
```

1 (a) - OLS Model

1 (b) - PLS Model

1 (c) - LASSO Model

1 (d) - Model Variants

1 (d, i) - PCR Model

Perform PCA analysis to see how Principal components explain variance

- Uses `numeric` data for Principal Component Analysis
- Then appends the `factor` data to the data *without NULL values*
- Finally, uses `stepAIC()` to best model data
- See interpretation at end

Get cleaned `numeric` and `factor` data frames

```
# After cleaning, two datasets that contain..

## Numeric data -----
hd.numericClean <- hd.Cleaned %>% select_if(is.numeric)

## Factors -----
hd.factorClean <- hd.Cleaned %>% dplyr::select(where(is.factor))

# Removing any columns with NA
removeColsWithNA <- function(df) {
  return( df[ , colSums(is.na(df)) == 0] )
}
hd.factorClean <- removeColsWithNA(hd.factorClean)

paste('Num. factor cols. removed due to null values:',
      ncol(hd.Cleaned %>% dplyr::select(where(is.factor))) - ncol(hd.factorClean) )
```

```
## [1] "Num. factor cols. removed due to null values: 16"
```

```
paste(ncol(hd.factorClean), 'factor cols. remain')
```

```
## [1] "22 factor cols. remain"
```

Perform PCA

```
# Principal component analysis on numeric data
pc.house <- prcomp(hd.numericClean %>% dplyr::select(-SalePrice), # do not include response var
                  center = TRUE, # Mean centered
                  scale = TRUE # Z-Score standardized
                  )

# See first 10 cumulative proportions
pc.house.summary <- summary(pc.house)
pc.house.summary$importance[, 1:10]
```

	PC1	PC2	PC3	PC4	PC5	PC6
## Standard deviation	2.655641	1.854044	1.61067	1.393107	1.170723	1.105003
## Proportion of Variance	0.227500	0.110890	0.08369	0.062600	0.044210	0.039390
## Cumulative Proportion	0.227500	0.338380	0.42207	0.484670	0.528890	0.568280

	PC7	PC8	PC9	PC10
## Standard deviation	1.061151	1.035538	1.006565	0.9997289
## Proportion of Variance	0.036320	0.034590	0.032680	0.0322400
## Cumulative Proportion	0.604600	0.639190	0.671870	0.7041100

Now we choose number of PC's that explain 75% of the variation

- Note this threshold is just a judgement call. No significance behind 75%

```
cumPropThreshold = 0.75 # The threshold

numPCs <- sum(pc.house.summary$importance['Cumulative Proportion', ] < cumPropThreshold)
paste0('There are ', numPCs, ' principal components that explain up to ', cumPropThreshold*100,
      '% of the variation in the data')
```

```
## [1] "There are 11 principal components that explain up to 75% of the variation in the data"
```

```
chosenPCs <- as.data.frame(pc.house$x[, 1:numPCs])
```

Join on the factor data

```
df.pcr <- cbind(SalePrice = hd.numericClean$SalePrice, chosenPCs, hd.factorClean)
```

Now, Apply predictions with PCR

- Linear model containing:
 - Principal components explaining 75% of variation in numeric data
 - Non-null factor data
 - *Predicted variable:* log(SalePrice)
- Then use `stepAIC()` to identify which variables are actually important for model

```
# Fit data using PC's, non-null factors
fit.pcr <- lm(log(SalePrice) ~ ., data = df.pcr)

# Reduce to only important variables
fit.pcrReduced <- stepAIC(fit.pcr, direction="both")
```

```
# View results
summary(fit.pcrReduced)
```

```
##
## Call:
## lm(formula = log(SalePrice) ~ PC1 + PC2 + PC3 + PC4 + PC5 + PC7 +
##      PC8 + PC10 + MSZoning + LandContour + LotConfig + Condition1 +
```

```

##      HouseStyle + RoofStyle + Exterior1st + ExterQual + ExterCond +
##      Foundation + CentralAir + KitchenQual + Functional + PavedDrive,
##      data = df.pcr)
##
## Residuals:
##      Min        1Q      Median        3Q        Max
## -0.66920 -0.05779  0.00279  0.06390  0.31705
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    11.789996   0.054155  217.710 < 2e-16 ***
## PC1              0.098465   0.002397   41.083 < 2e-16 ***
## PC2              0.008533   0.003243    2.631 0.008646 **
## PC3             -0.053870   0.003278  -16.435 < 2e-16 ***
## PC4             -0.018160   0.003594   -5.053 5.21e-07 ***
## PC5              0.050661   0.003387   14.958 < 2e-16 ***
## PC7              0.007044   0.003362    2.095 0.036427 *
## PC8             -0.016103   0.003487   -4.618 4.40e-06 ***
## PC10             0.011869   0.003541    3.352 0.000834 ***
## MSZoningRH      -0.058055   0.039955   -1.453 0.146548
## MSZoningRL      -0.031020   0.020263   -1.531 0.126128
## MSZoningRM      -0.110454   0.021970   -5.028 5.94e-07 ***
## LandContourHLS    0.081817   0.026979    3.033 0.002491 **
## LandContourLow    0.001438   0.028780    0.050 0.960167
## LandContourLvl   -0.006761   0.018252   -0.370 0.711128
## LotConfigCulDSac  0.046494   0.015183    3.062 0.002259 **
## LotConfigInside   0.006374   0.009135    0.698 0.485496
## LotConfigother   -0.005573   0.019309   -0.289 0.772931
## Condition1Feedr   0.054752   0.024923    2.197 0.028273 *
## Condition1Norm    0.093747   0.020579    4.555 5.91e-06 ***
## Condition1RR      0.056079   0.029465    1.903 0.057312 .
## Condition1Other   0.018423   0.031438    0.586 0.557998
## HouseStyle1Story -0.066780   0.015570   -4.289 1.98e-05 ***
## HouseStyle2Story -0.017051   0.015214   -1.121 0.262706
## HouseStyleSLvl   -0.033535   0.020570   -1.630 0.103365
## HouseStyleOther  -0.055545   0.020119   -2.761 0.005875 **
## RoofStyleHip      0.015341   0.009404    1.631 0.103149
## RoofStyleother    0.102425   0.024857    4.121 4.11e-05 ***
## Exterior1stMetalSd 0.028893   0.012682    2.278 0.022934 *
## Exterior1stVinylSd 0.027004   0.011415    2.366 0.018198 *
## Exterior1stWd Sdng -0.006310   0.013395   -0.471 0.637717
## Exterior1stOther   0.036398   0.011499    3.165 0.001598 **
## ExterQualAvg      -0.040920   0.011707   -3.495 0.000495 ***
## ExterQualBelowAvg -0.116457   0.046627   -2.498 0.012670 *
## ExterCondAvg       0.020386   0.011733    1.738 0.082620 .
## ExterCondBelowAvg -0.007850   0.032377   -0.242 0.808472
## FoundationCBlock   0.001430   0.014332    0.100 0.920544
## Foundationother    0.024638   0.025685    0.959 0.337685
## FoundationPConc    0.051223   0.016776    3.053 0.002326 **
## CentralAirY       0.049710   0.016954    2.932 0.003449 **
## KitchenQualAvg    -0.021347   0.010486   -2.036 0.042052 *
## KitchenQualBelowAvg -0.042348   0.025577   -1.656 0.098109 .
## FunctionalMaj2    -0.208427   0.062071   -3.358 0.000817 ***
## FunctionalMin1     0.032238   0.039058    0.825 0.409361

```

```
## FunctionalMin2      0.021286    0.038067    0.559 0.576172
## FunctionalMod       0.001628    0.044465    0.037 0.970795
## FunctionalTyp       0.097215    0.031899    3.048 0.002371 **
## PavedDriveP        -0.008323    0.025418   -0.327 0.743401
## PavedDriveY         0.049301    0.016297    3.025 0.002551 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1049 on 951 degrees of freedom
## Multiple R-squared:  0.9205, Adjusted R-squared:  0.9165
## F-statistic: 229.5 on 48 and 951 DF,  p-value: < 2.2e-16
```

```
# Key diagnostics
est.pcr <- summary(fit.pcrReduced)

# Get the RMSE and R Squared of the model
keyDiagnostics.pcr <- data.frame(Model    = 'PCR',
                                Notes     = 'lm',
                                Hyperparameters = 'N/A',
                                RMSE      = sqrt(mean(fit.pcrReduced$residuals^2)),
                                Rsquared   = est.pcr$adj.r.squared)

# Show output
keyDiagnostics.pcr %>% kable()
```

Model	Notes	Hyperparameters	RMSE	Rsquared
PCR	lm	N/A	0.1023139	0.9165338

Interpretation of PCR Model

Please note all interpretations below are approximate, given the `stepAIC()` uses stochastic modeling.

Model performance evaluation:

- See that around 28 of the variables cannot be explained by random chance, with a probability of 90% or more (see significance codes above)
- Standard errors range from \pm 1-5%, with average around 2%. Larger values may indicate higher uncertainty of the estimated coefficients.
- This model explains around 92% of the variation in the `log(SalePrice)`. See Adjusted R-Squared for reference.
- Note this model may exhibit selection bias, since the data excludes factor data with null values in the variable.
- This model would likely do well for prediction of `log(SalePrice)`, given the small range of standard errors, high adjusted R squared, and number of significant variables. This model would obviously not do well for inference, given we are using principal components that mask the numeric data.

Practical significance evaluation:

- The principal components contribute positively about 20% of the sale price of the home

- Residential Medium Density (**MSZoningRM**) reduces the home price by around 12%, with a standard error of around 2%.
- If the exterior quality is below average (**ExterQualBelowAvg**), it reduces the home price by around 12%, with a standard error of around 5%.
- If the functionality of the home has 2 major deductions (**FunctionalMaj2**), it reduces the home price by around 20%, with a standard error of around 6%. While having typical functionality (**FunctionalTyp**) increases the home sale price by nearly 10%, with a standard error of 3%.
- See other coefficients of the data for other variables.

Visualizatoion of PCR Model - Predicted vs. Actuals

Function to compare predicted vs. observed values

```
# Function to compare predicted vs. actual (observed) regression outputs
predictedVsObserved <- function(predicted, observed, modelName, outcomeName = 'Log(SalePrice)') {

  ## Create dataset for predicted vs. actuals
  comparison <- data.frame(observed = observed,
                           predicted = predicted) %>%

  # Row index
  mutate(ID = row_number()) %>%

  # Put in single column
  pivot_longer(cols = c('observed', 'predicted'),
               names_to = 'metric',
               values_to = 'value')

  # Plot --- Observed vs. Actuals across all variables in data
  variationScatter <- comparison %>%
    ggplot(aes(x = ID,
               y = value,
               color = metric
               )
           ) +
    geom_point(alpha = 0.5, size = 1) +

    labs(title = 'Variation in Predicted vs. Observed Data',
         subtitle = paste('Model:', modelName),
         x = 'X', y = outcomeName) +
    theme_minimal() + theme(legend.title = element_blank(),
                           legend.position = 'top') +
    scale_color_manual(values = c('grey60', 'palegreen3'))

  print(variationScatter)

  # Limit for x and y axis for scatter of predicted vs. observed
  axisLim = c( min(c(predicted, observed)), max(c(predicted, observed)) )

  # Simple comparison of data
  plot(x = observed,
       y = predicted,
       main = 'PCR Model - Actual (Observed) vs. Predicted\n',
       xlab = paste('Observed Values -', outcomeName),
       ylab = paste('Predicted Values -', outcomeName),
       pch = 16,
       cex = 0.75,
       col = alpha('steelblue3', 1/4),
       xlim = axisLim,
       ylim = axisLim)
```

```

)

# Add the Predicted vs. actual line
abline(lm(predicted ~ observed), col = 'steelblue3', lwd = 2)
mtext('Predicted ~ Actual', side = 3, adj = 1, col = 'steelblue3')

# Add line for perfectly fit model
abline(0,1, col = alpha('tomato3', 0.8), lwd = 2)
mtext('Perfectly Fit Model', side = 1, adj = 0, col = 'tomato3')
}

```

View results of the PCR Model

- See that the variation in the data is very closely resembled actual by changes in independent variables
- Implication? This model fits its own data well, but what is not know if it can predict out of sample data.
- Note that it the data (blue) deviates slightly from perfect line model (red), indicating that the model is slightly skewed from predicted and actual data.


```

# How do the two compare?
predictedVsObserved(observed = log(df.pcr$SalePrice),
                    predicted = predict(fit.pcrReduced),
                    modelName = 'PCR')


```

Variation in Predicted vs. Observed Data

Model: PCR



PCR Model – Actual (Observed) vs. Predicted



1 (d, ii) - SVR Model

Model Setup

```
ctrl <- trainControl(method = "repeatedcv",
                     number = 5, # 5 fold cross validation
                     repeats = 2 # 5 repeats
                     )

# The data (pmm imputed values, and only whole columns without NA. Does not omit outliers)
df.svm <- cbind(SalePrice = hd.numericClean$SalePrice,
               hd.numericClean, hd.factorClean)
```

Tune Model

```
# Train and tune the SVM - first broad tuning on cost only
svm.tune <- train(data = df.svm,
                 log(SalePrice) ~ .,
                 method = "svmRadial",          # Radial kernel
                 tuneLength = 9,                # 9 values of the cost function
                 preProc = c("center","scale"), # Center and scale data
                 trControl = ctrl)
```

Diagnostics of Model

- Note all numbers mentioned below are approximate
- See that the R Squared of the model is around 0.86, and RMSE is 0.14
- See that the model predicts the data well.
- Also, note that the model predicts the data with less error than the linear model. See this from the RMSE or scatter plot of predicted values.

```
# Get the RMSE and R Squared of the model
hyperparameters.svm = list('C' = svm.tune[["finalModel"]@param[["C"]],
                          'Epsilon' = svm.tune[["finalModel"]@param[["epsilon"]])

keyDiagnostics.svm <- data.frame(Model = 'SVM',
                                Notes = 'caret and svmRadial',
                                Hyperparameters = paste('C =', hyperparameters.svm$C, ',',
                                                         'Epsilon =', hyperparameters.svm$Epsilon)
                                )

keyDiagnostics.svm <- cbind(keyDiagnostics.svm,
                          svm.tune$results %>%
                            filter(C == hyperparameters.svm$C) %>%
                            dplyr::select(RMSE, Rsquared)
                          )

# Show output
keyDiagnostics.svm %>% kable()
```

Model	Notes	Hyperparameters	RMSE	Rsquared
SVM	caret and svmRadial	C = 4 , Epsilon = 0.1	0.1419626	0.8521295


```
# Final model?
svm.tune$finalModel
```

```
## Support Vector Machine object of class "ksvm"
##
## SV type: eps-svr (regression)
## parameter : epsilon = 0.1 cost C = 4
##
## Gaussian Radial Basis kernel function.
## Hyperparameter : sigma = 0.00772251903100742
##
## Number of Support Vectors : 661
##
## Objective Function Value : -161.4761
## Training error : 0.012257
```

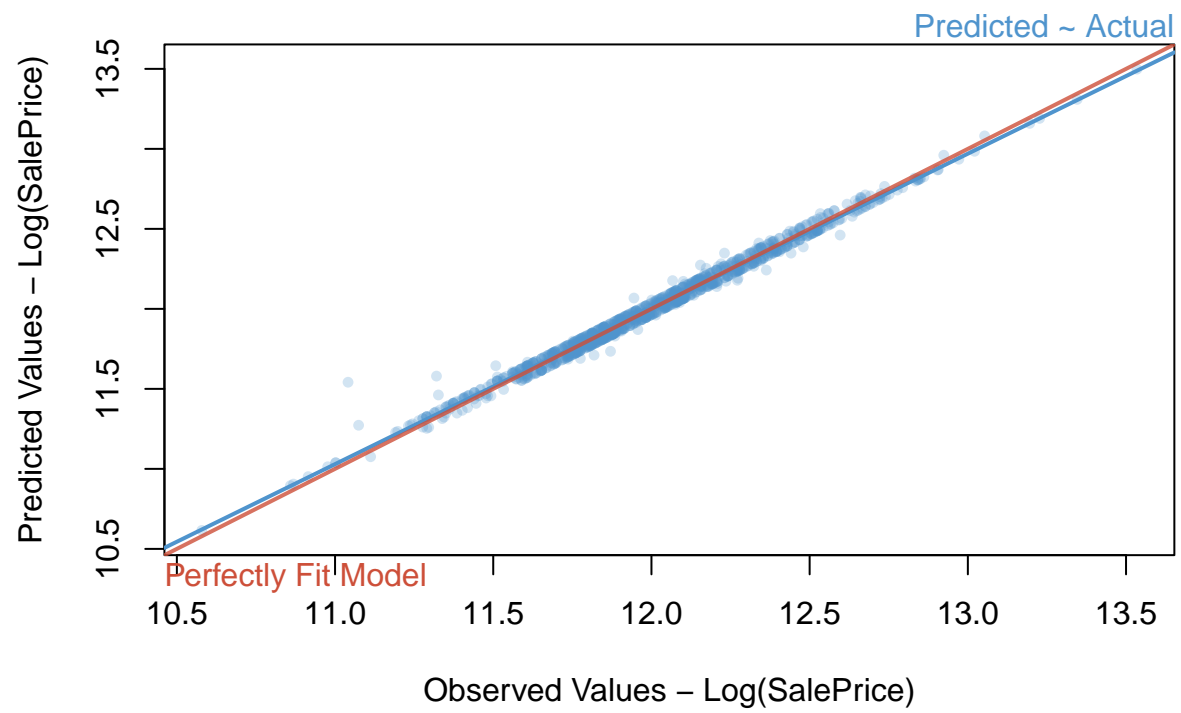
```
# How do the two compare?
predictedVsObserved(
  observed = log(df.svm$SalePrice),
  predicted = predict(svm.tune, df.svm),
  modelName = 'SVM')
```

Variation in Predicted vs. Observed Data

Model: SVM



PCR Model – Actual (Observed) vs. Predicted



1 (d, iii) - MARS Model

Summary Table of Model Performance

```
rbind(keyDiagnostics.pcr, keyDiagnostics.svm) %>% kable()
```

Model	Notes	Hyperparameters	RMSE	Rsquared
PCR	lm	N/A	0.1023139	0.9165338
SVM	caret and svmRadial	C = 4 , Epsilon = 0.1	0.1419626	0.8521295