

ISE 5103 Intelligent Data Analytics

Homework 5 - Modeling

Daniel Carpenter & Sonaxy Mohanty

October 2022

Contents

Packages	3
General Data Prep	3
Read Data	3
Impute Missing Values with PMM	4
Factor Level Collapse - Create Other Bin for Columns over 4 Unique Values	5
Remove Outliers from Numeric Data	6
Exploratory Data Analysis	7
Checking the distribution of Sale Price of houses	7
Correlation between features in the dataset	7
1 (a) - OLS Model	9
i.	9
Hold-out validation set	9
Fit the OLS Model	9
Fit the Model	9
ii. Complete analysis of the residuals	12
1 (b) - PLS Model	16
Model Setup	16
1 (c) - LASSO Model	19
Model Setup	19
Fit the Model	19
1 (d) - Model Variants	21
1 (d, i) - PCR Model	21
Model Setup	21
Fit the Model	21
View and Interpret Results	23
1 (d, ii) - SVR Model	27
Model Setup	27
Fit the Model	27
View and Interpret Results	27
1 (d, iii) - MARS Model	29
Fit the Model	29
View and Interpret Results	29
Summary Table of Model Performance	31

Packages

```
# Data Wrangling
library(tidyverse)

# Modeling
library(MASS)
library(caret) # Modeling variants like SVM
library(earth) # Modeling with Mars
library(pls) #Modeling with PLS
library(glmnet) #Modeling with LASSO

# Aesthetics
library(knitr)
library(cowplot) # multiple ggplots on one plot with plot_grid()
library(scales)
library(kableExtra)
library(ggplot2)

#Hold-out Validation
library(caTools)

#Data Correlation
library(GGally)
library(regclass)

#RMSE Calculation
library(Metrics)

#p-value for OLS model
library(broom)

#ncvTest
library(car)
```

General Data Prep

Read Data

```
# Convert all character data to factor
hd <- read.csv('housingData.csv', stringsAsFactors = TRUE) %>%

# creates new variables age, ageSinceRemodel, and ageofGarage and
dplyr::mutate(age = YrSold - YearBuilt,
              ageSinceRemodel = YrSold - YearRemodAdd,
              ageofGarage = ifelse(is.na(GarageYrBlt), age, YrSold - GarageYrBlt)) %>%

# remove some columns used in the above calculations
dplyr::select(!c(Id,YrSold ,
                 MoSold, YearBuilt, YearRemodAdd))

#str(hd)
```

Impute Missing Values with PMM

Make data set of `numeric` variables

Make data set of `factor` variables

For each column with missing data, impute missing values with PMM

- Done with function `imputeWithPMM()` function
- Applies function via `dplyr` logic
- Note `seeImputation()` function to visualize the imputation from prior homework 4, not shown for simplicity in viewing

Create function to impute via PMM

Apply PMM function to numeric data containing null values

```
## [1] "LotFrontage" "MasVnrArea" "GarageYrBlt"
```

```
## [1] "For imputation results of LotFrontage, see OutputPMM/Imputation_With_PMM_LotFrontage.pdf"
```

```
## [1] "For imputation results of MasVnrArea, see OutputPMM/Imputation_With_PMM_MasVnrArea.pdf"
```

```
## [1] "For imputation results of GarageYrBlt, see OutputPMM/Imputation_With_PMM_GarageYrBlt.pdf"
```

Factor Level Collapse - Create Other Bin for Columns over 4 Unique Values

```
## [1] "Before cleaning, there are 14 factor columns with more than 4 unique values"
```

```
## [1] "After cleaning, there are 14 columns with more than 4 unique values (omitting NA's)"
```

Remove Outliers from Numeric Data

- Since there are so many outliers, we are only going to remove some outliers
- If you count the number of outliers by column, the 75% of columns contain less than 50 outliers.
- However, some contain up to 200. Since remove ALL outliers would reduce the size of the data to less than 300 observations, we are removing up to 50 per column.

```
## [1] "Of the columns with outliers, removed up to 75th percentile of num. outliers."  
## [1] "See that the 75th percentile of columns with outliers contain 51.75 outliers"
```

Exploratory Data Analysis

Checking the distribution of Sale Price of houses

```
hist(hd.CleanedNoOutliers$SalePrice,  
     col = 'skyblue4',  
     main = 'Distribution of Sale Price of houses',  
     xlab = 'House Price')
```



- After removing the desired outliers, we can see that the distribution of Sale Price looks like a normal distribution with few outliers on the right tail.

Correlation between features in the dataset

```
ggcorr(hd.CleanedNoOutliers, geom='blank', label=T, label_size=3, hjust=1,  
       size=3, layout.exp=2) +  
  geom_point(size = 4, aes(color = coefficient > 0, alpha = abs(coefficient) >= 0.5)) +  
  scale_alpha_manual(values = c("TRUE" = 0.25, "FALSE" = 0)) +  
  guides(color = F, alpha = F)
```


1 (a) - OLS Model

i.

Hold-out validation set

- Since, we have deleted some of the outlier values during data pre-processing, using 10% of the data as test and remaining 90% as train

```
idx <- sample(nrow(hd.CleanedNoOutliers), nrow(hd.CleanedNoOutliers)*0.1)
test <- hd.CleanedNoOutliers[idx,]
train <- hd.CleanedNoOutliers[-idx,]
```

Fit the OLS Model

Model 1:

- Linear model containing:
 - *Independent variables:* GarageArea + GarageCars + TotRmsAbvGrd + FullBath + GrLivArea + X1stFlrSF + TotalBsmtSF + OverallQual
 - *Predicted variable:* SalePrice

```
ols.mdl1 <- lm(SalePrice ~ GarageArea + GarageCars + TotRmsAbvGrd
               + FullBath + GrLivArea + X1stFlrSF + TotalBsmtSF + OverallQual, data=train)
```

- **For Model 1:** Adjusted R-squared is 0.821, AIC is 16695.31 and BIC is 16741.28 and RMSE is 20541.22.
- Still trying to improve the existing model.
- No multicollinearity detected.

Model 2:

- This model created is based on **Principal Component Analysis**.
 - Uses **numeric** data for Principal Component Analysis
 - Then appends the **factor** data to the data *without NULL values*
 - Finally, uses **stepAIC()** to best model data

Now we choose number of PC's that explain 75% of the variation

- Note this threshold is just a judgement call. No significance behind 75%

```
## [1] "There are 9 principal components that explain up to 75% of the variation in the data"
```

Fit the Model

- Linear model containing:
 - Principal components explaining 75% of variation in **numeric** data
 - Non-null **factor** data
 - *Predicted variable:* SalePrice
- Then use **stepAIC()** to identify which variables are actually important for model

```
# Fit data using PC's, non-null factors
fit.ols <- lm(SalePrice ~ ., data = df.ols)
```

```
# Reduce to only important variables
ols.mdl2 <- stepAIC(fit.ols, direction="both")
```

- Reporting all the variables of the best model (Model 2):

Coefficient estimates:

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	132084.03696	10595.8674	12.46561817	2.991747e-32
## PC1	12562.26171	475.1182	26.44028795	2.507707e-106
## PC3	-6602.94706	730.4852	-9.03912523	1.644606e-18
## PC4	-3506.79243	521.5626	-6.72362733	3.774185e-11
## PC5	1642.77240	660.5434	2.48700148	1.312263e-02
## PC6	-7099.91044	708.8347	-10.01631302	4.110301e-22
## PC7	-4157.93100	641.9053	-6.47748309	1.795685e-10
## PC8	-1289.27140	630.8678	-2.04364756	4.137549e-02
## PC9	2488.81664	648.7043	3.83659662	1.364335e-04
## MSZoningRH	-16613.40293	8836.3403	-1.88012258	6.052115e-02
## MSZoningRL	-5258.83266	3824.4485	-1.37505648	1.695695e-01
## MSZoningRM	-13506.46541	4230.2960	-3.19279438	1.474376e-03
## LandContourHLS	13972.61581	5334.1867	2.61944633	9.004838e-03
## LandContourLow	9540.10875	5735.3207	1.66339587	9.669671e-02
## LandContourLvl	64.73445	3467.5304	0.01866875	9.851109e-01
## LotConfigCulDSac	4274.74389	2844.1948	1.50297153	1.333134e-01
## LotConfigInside	-1838.95995	1765.6924	-1.04149507	2.980182e-01
## LotConfigOther	-5876.77666	3650.4997	-1.60985542	1.078963e-01
## NeighborhoodNames	-7654.95043	2813.1977	-2.72108512	6.674244e-03
## NeighborhoodOldTown	-3948.31147	3920.4195	-1.00711454	3.142401e-01
## NeighborhoodOther	-2980.34849	3420.4132	-0.87134165	3.838770e-01
## NeighborhoodOther	-797.10902	2218.9501	-0.35922801	7.195367e-01
## Condition1Feedr	2878.23898	4982.7025	0.57764616	5.636954e-01
## Condition1Norm	10734.21247	4130.6829	2.59865325	9.563062e-03
## Condition1RR	-322.85461	5575.1128	-0.05790997	9.538375e-01
## Condition1Other	1953.48688	6811.7895	0.28678028	7.743684e-01
## BldgType2fmCon	-14195.57637	7464.0298	-1.90186491	5.761364e-02
## BldgTypeDuplex	-6658.56135	8781.2982	-0.75826617	4.485559e-01
## BldgTypeTwnhs	-13769.29172	4896.2727	-2.81219871	5.063128e-03
## BldgTypeTwnhsE	-3545.85572	3923.0796	-0.90384495	3.663996e-01
## HouseStyle1Story	-4779.78828	2627.0620	-1.81944250	6.928615e-02
## HouseStyle2Story	4416.44039	2895.5935	1.52522804	1.276699e-01
## HouseStyleSLvl	-3546.26247	3901.8213	-0.90887361	3.637408e-01
## HouseStyleOther	-4007.69457	3776.4927	-1.06122131	2.889683e-01
## RoofStyleHip	2937.37319	1790.3806	1.64064180	1.013370e-01
## RoofStyleOther	24330.40171	5534.2609	4.39632361	1.278419e-05
## Exterior1stMetalSd	8454.65291	7172.9232	1.17869000	2.389364e-01
## Exterior1stVinylSd	-2276.67820	7385.5030	-0.30826312	7.579771e-01
## Exterior1stWd Sdng	-6159.99005	5194.6354	-1.18583685	2.361034e-01
## Exterior1stOther	9225.20645	3671.2218	2.51284366	1.220784e-02
## Exterior2ndMetalSd	-914.74166	7290.6898	-0.12546709	9.001911e-01
## Exterior2ndVinylSd	9503.71922	7542.2462	1.26006483	2.080808e-01
## Exterior2ndWd Sdng	11037.76569	5309.4984	2.07887164	3.800568e-02
## Exterior2ndOther	-3442.93039	3631.0054	-0.94820305	3.433649e-01
## ExterQualAvg	-10817.79756	2236.4785	-4.83697814	1.633866e-06
## ExterQualBelowAvg	27451.15724	11846.0210	2.31733148	2.078350e-02
## ExterCondAvg	4706.41525	2178.2433	2.16064721	3.107323e-02

```
## ExterCondBelowAvg      683.82765  6425.4413   0.10642501  9.152767e-01
## KitchenQualAvg        -4410.58396  1906.6264  -2.31329216  2.100578e-02
## KitchenQualBelowAvg   1613.53530  4774.9596   0.33791601  7.355314e-01
## FunctionalMaj2        -19264.90742 14110.3202  -1.36530618  1.726110e-01
## FunctionalMin1         15263.09959  7587.4961   2.01161218  4.465720e-02
## FunctionalMin2         18176.27161  7650.9839   2.37567768  1.779456e-02
## FunctionalMod          27704.34340 12113.3971   2.28708291  2.249916e-02
## FunctionalTyp          28660.70584  6454.1324   4.44067525  1.047319e-05
## PavedDriveP           -1166.21480  5071.0772  -0.22997378  8.181817e-01
## PavedDriveY            7233.65295  3292.1203   2.19726267  2.834024e-02
```

p-values:

```
##          value
## 2.12313e-290
```

Adjusted R-squared:

```
## [1] 0.8849292
```

AIC:

```
## [1] 16381.04
```

BIC:

```
## [1] 16647.67
```

VIF:

```
##          GVIF Df GVIF^(1/(2*Df))
## PC1          4.080965  1      2.020140
## PC3          4.346983  1      2.084942
## PC4          1.469448  1      1.212208
## PC5          1.582655  1      1.258036
## PC6          1.684946  1      1.298055
## PC7          1.237646  1      1.112495
## PC8          1.148415  1      1.071641
## PC9          1.158146  1      1.076172
## MSZoning      3.645275  3      1.240571
## LandContour   1.560030  3      1.076933
## LotConfig     1.352988  3      1.051677
## Neighborhood  5.886935  4      1.248062
## Condition1    1.644309  4      1.064138
## BldgType      6.380121  4      1.260676
## HouseStyle    5.536189  4      1.238515
## RoofStyle     1.427511  2      1.093062
## Exterior1st   6632.913003  4      3.004091
## Exterior2nd   6515.077892  4      2.997367
## ExterQual      4.218565  2      1.433148
## ExterCond      1.614037  2      1.127141
## KitchenQual    2.825011  2      1.296448
## Functional     2.105329  5      1.077288
## PavedDrive     1.600230  2      1.124723
```

RMSE:

```
## [1] 15929.13
```

- So, we can say that using PCA followed by stepAIC the OLS regression model is better as compared to the other OLS model built.
- There is also no multicollinearity found in the model as the VIF values are less than 10.

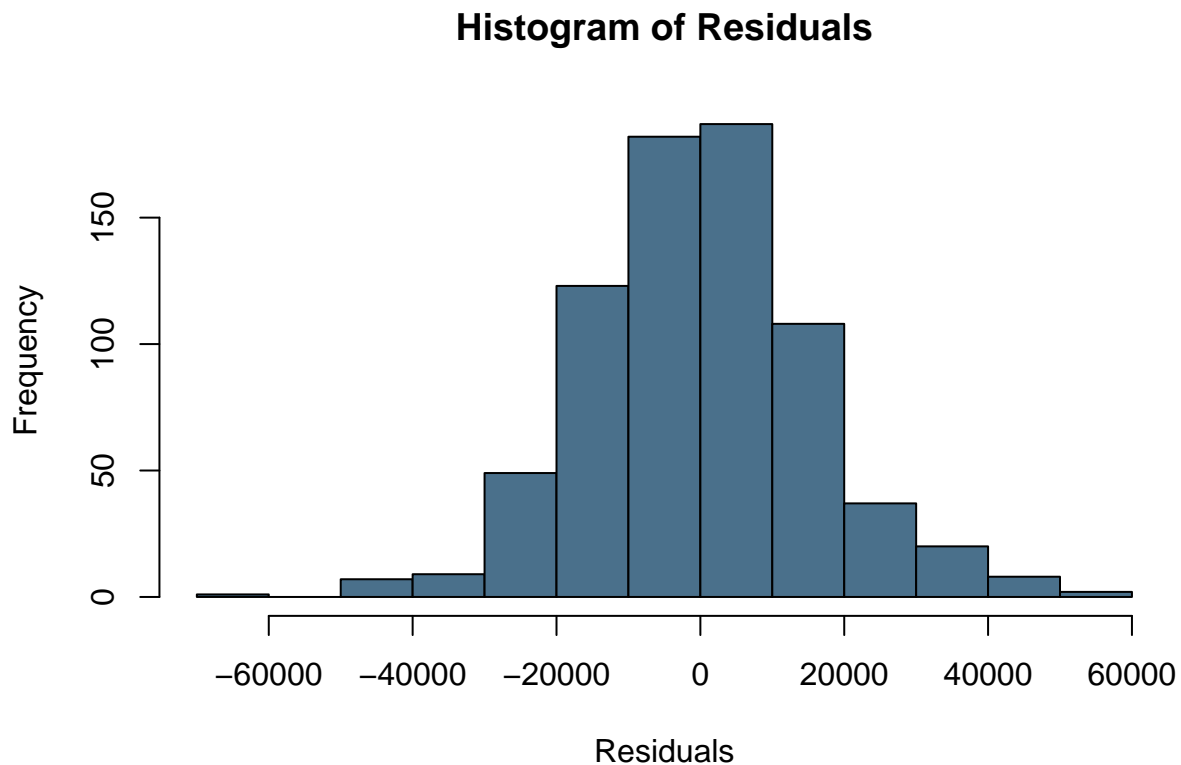
Model	Notes	Hyperparameters	RMSE	Rsquared
OLS	lm + 2-way interactions	N/A	15929.13	0.8849292

ii. Complete analysis of the residuals

A linear regression model is considered fit if the below assumptions are met:

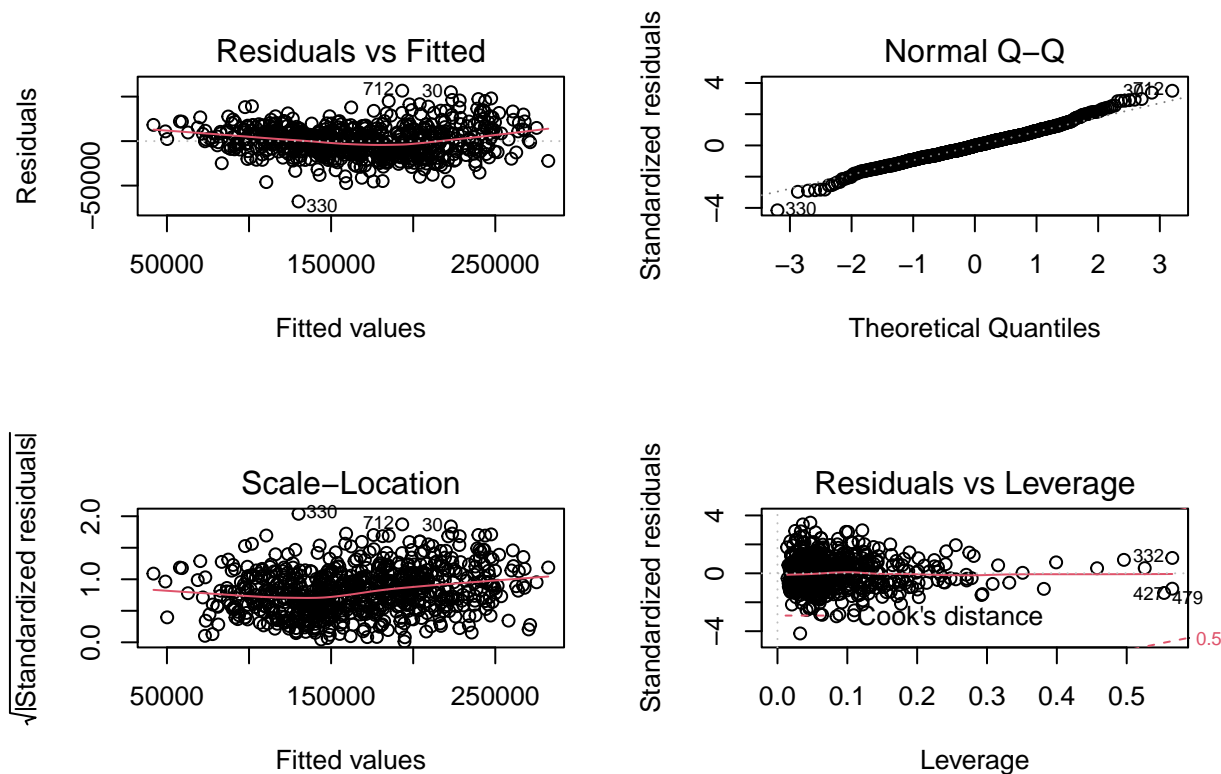
- **Residuals should follow normal distribution**
- **There should be no heteroscedasticity**
- **There should be no multicollinearity**

```
hist(ols.mdl2$residuals,
     col = 'skyblue4',
     main = 'Histogram of Residuals',
     xlab = 'Residuals')
```



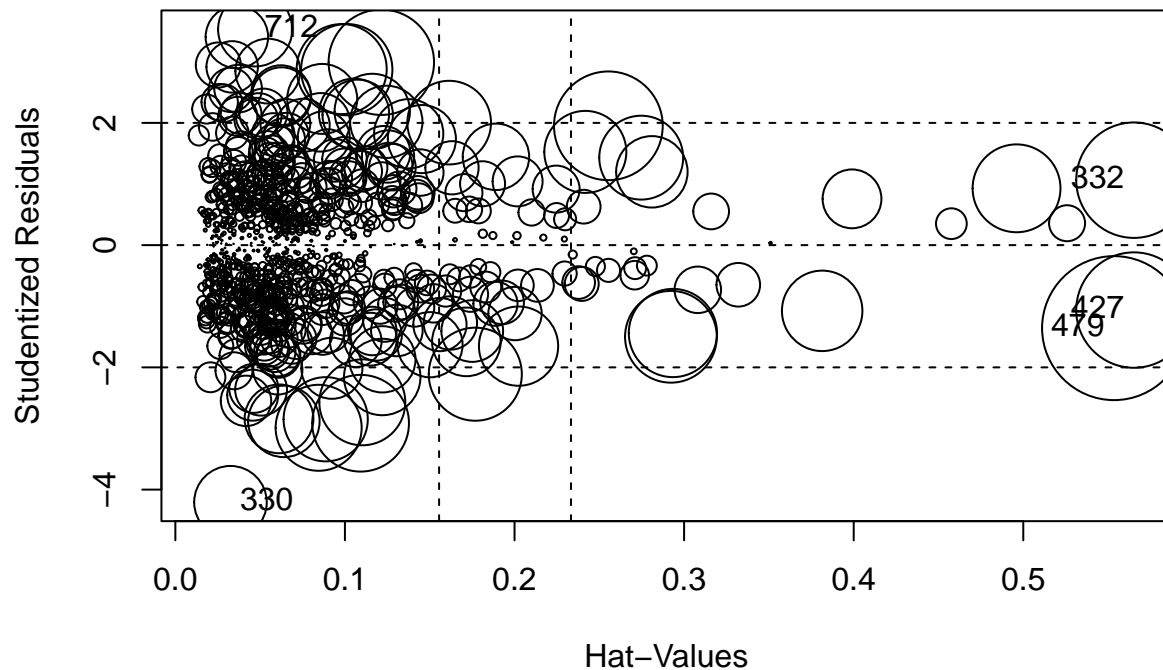
We can see that the residuals are normally distributed.

```
par(mfrow=c(2,2)) #combining multiple plots together
plot(ols.mdl2)
```



- From the *Residuals vs Fitted* plot, we can see there are points above and below the 0 line.
- There is also a pattern seen like a **slight curvature pattern** which indicates that there may be a systematic lack of fit.
- From the *Normal Q-Q* plot, we can see that most of the points are **very close to the dotted line**, indicating that the residuals follow a normal distribution, except some points which might be outliers which maybe affecting the regression line fit of data.
- Here the *Scale-Location* plot suggests that the red line is roughly horizontal across the plot and the spread of magnitude looks unequal, at some fitted values there are more residuals as compared to other like the ones in between 120000 and 210000, indicating some heteroskedasticity.
- From the *Residuals vs Leverage* plot, we can see that there are no influential points in our regression model. We need to check `influencePlot` to see if we are missing any leverage.

```
influencePlot(ols.mdl2)
```



```
##      StudRes      Hat      CookD
## 330 -4.206583 0.03250006 0.01017704
## 332  1.062619 0.56525366 0.02575165
## 427 -1.062619 0.56525366 0.02575165
## 479 -1.357941 0.55353100 0.04005859
## 712  3.530266 0.04707287 0.01062059
```

- We can now see some high influential points for the fitted values.

```
#ncv Test
ncvTest(ols.mdl2)
```

```
## Non-constant Variance Score Test
## Variance formula: ~ fitted.values
## Chisquare = 45.09124, Df = 1, p = 1.8806e-11
```

Since p-value is less than significance level (α) of 0.05, that means we reject the null hypothesis of constant error variance which indicates heteroscedasticity.

```
VIF(ols.mdl2)
```

```
##      GVIF Df GVIF^(1/(2*Df))
## PC1      4.080965 1      2.020140
## PC3      4.346983 1      2.084942
## PC4      1.469448 1      1.212208
## PC5      1.582655 1      1.258036
## PC6      1.684946 1      1.298055
## PC7      1.237646 1      1.112495
## PC8      1.148415 1      1.071641
```

## PC9	1.158146	1	1.076172
## MSZoning	3.645275	3	1.240571
## LandContour	1.560030	3	1.076933
## LotConfig	1.352988	3	1.051677
## Neighborhood	5.886935	4	1.248062
## Condition1	1.644309	4	1.064138
## BldgType	6.380121	4	1.260676
## HouseStyle	5.536189	4	1.238515
## RoofStyle	1.427511	2	1.093062
## Exterior1st	6632.913003	4	3.004091
## Exterior2nd	6515.077892	4	2.997367
## ExterQual	4.218565	2	1.433148
## ExterCond	1.614037	2	1.127141
## KitchenQual	2.825011	2	1.296448
## Functional	2.105329	5	1.077288
## PavedDrive	1.600230	2	1.124723

Generally, VIF values which are greater than 5 or 7 are the cause of multicollinearity which we do not see in our model.

Improving the current model:

* To improve our model, we need to remove some influential observations from our model and then fit the regression model to the data.

* We can re-build the model with new predictors.

* We can also perform variable transformation such as Box-Cox or use better evolved models like SVR, PCR etc., and see how it works.

1 (b) - PLS Model

Model Setup

- Using the whole data set after PMM imputation and factor level collapsing without omitting any outliers
- Using the predictors - GarageArea, GarageCars, TotRmsAbvGrd, FullBath, GrLivArea, X1stFlrSF, TotalBsmtSF, OverallQual which has strong correlations with response variable - SalePrice

```
#creating a PLS model to predict the log of the sale price
#using 5-fold CV
```

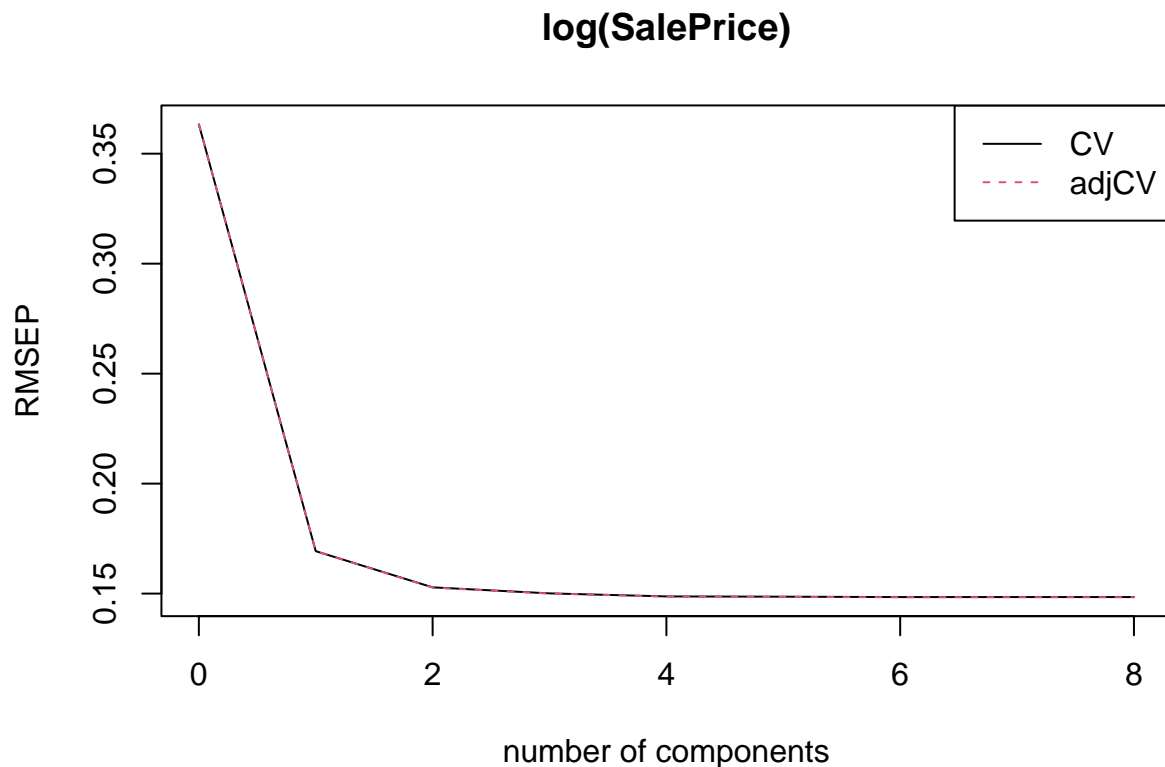
```
pls.model <- plsr(log(SalePrice) ~ GarageArea + GarageCars + TotRmsAbvGrd
  + FullBath + GrLivArea + X1stFlrSF + TotalBsmtSF + OverallQual,
  data=hd.Cleaned, scale=TRUE, validation='CV', k=5)
```

- Hyperparameter tuning to determine the number of PLS components with RMSE as the error metric

```
#report chart
summary(pls.model)
```

```
## Data:      X dimension: 1000 8
## Y dimension: 1000 1
## Fit method: kernelpls
## Number of components considered: 8
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           0.3633  0.1693  0.1528  0.1501  0.1487  0.1486  0.1484
## adjCV        0.3633  0.1693  0.1527  0.1500  0.1487  0.1485  0.1483
##      7 comps  8 comps
## CV          0.1484  0.1484
## adjCV       0.1484  0.1484
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           54.34  62.66  74.93  79.61  83.32  95.82  97.73
## log(SalePrice) 78.36  82.58  83.18  83.49  83.60  83.60  83.60
##      8 comps
## X           100.0
## log(SalePrice) 83.6
```

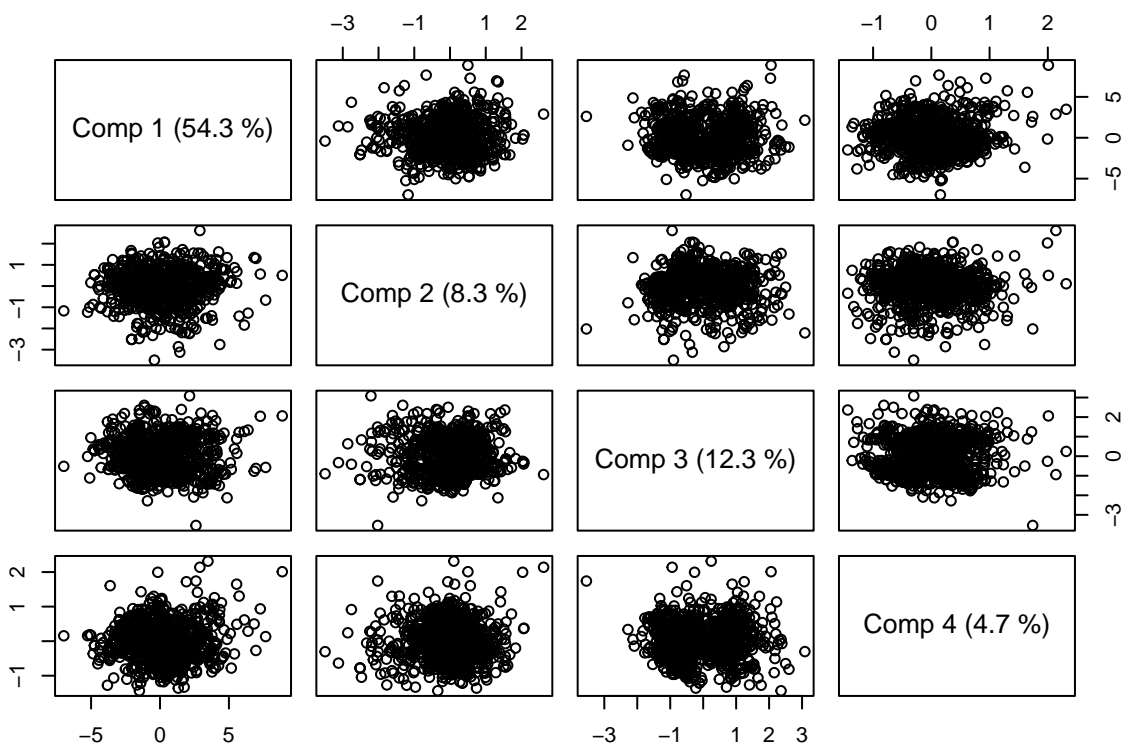
```
plot(RMSEP(pls.model), legendpos="topright")
```

- From the table, we can see that if we use 6 PLS components only in our model, the RMSE drops to 0.1486 and after that even if we keep adding components the RMSE still is the same.
- Though we are eyeballing the CV component, but from the plot we can see that fitting 4 PLS components is enough because even if we are adding 2 more components there is not much difference in the CV component.
- Using the final model with four PLS components to make predictions

```
final.pls <- plsr(log(SalePrice) ~ GarageArea + GarageCars + TotRmsAbvGrd
  + FullBath + GrLivArea + X1stFlrSF + TotalBsmtSF + OverallQual,4,
  data=hd.Cleaned, scale=TRUE, validation='CV', k=5)

plot(final.pls, plotype = "scores", comps = 1:4)
```



- From the above plot, we can see that by using only four PLS components we can describe about 80% of the variation in the response variable.
- Metric Calculations:

Model	Notes	Hyperparameters	RMSE	Rsquared
PLS	pls	ncomp = 4	0.1474771	0.0218368

- If we now compare between our preferred OLS model and PLS model on basis of RMSE values, we can see that PLS model's efficiency is much higher.
- RMSE for chosen OLS model was `ols.mdl2.rsme` whereas for PLS model is `0.1475`.

1 (c) - LASSO Model

Model Setup

- We first setup our cross-validation strategy
- Then create a dataframe with PMM imputed values, and only whole columns without NA. Does not omit outliers
- Then we train the model using `glmnet` which actually fits the elastic net

```
ctrl <- trainControl(method = "repeatedcv",
                     number = 5, # 5 fold cross validation
                     repeats = 2 # 2 repeats
                     )

# The data (PMM imputed values, and only whole columns without NA. Does not omit outliers)
df.lasso <- cbind(SalePrice = hd.numericClean$SalePrice,
                  hd.numericClean, hd.factorClean)
```

Fit the Model

```
# Train and tune the SVM
fit.lasso <- train(data = df.lasso,
                  log(SalePrice) ~ .,
                  method = "glmnet", # Elastic net
                  preProc = c("center", "scale"), # Center and scale data
                  tuneLength = 10, #10 values of alpha and 10 lamda values for each
                  trControl = ctrl)
```

- The variables with non-zero coefficients of the final model:

```
lasso.coeff <- drop(coef(fit.lasso$finalModel, fit.lasso$bestTune$lambda))

lasso.coeff[lasso.coeff != 0]
```

##	(Intercept)	MSSubClass	LotFrontage	LotArea
##	1.200247e+01	-7.152110e-05	3.440567e-03	2.420300e-02
##	OverallQual	OverallCond	MasVnrArea	BsmtFinSF1
##	7.994212e-02	4.961582e-02	2.014865e-03	3.054406e-02
##	BsmtFinSF2	TotalBsmtSF	LowQualFinSF	GrLivArea
##	3.951459e-03	4.266567e-02	-1.226020e-03	1.330731e-01
##	BsmtFullBath	FullBath	HalfBath	BedroomAbvGr
##	1.013981e-02	6.445650e-05	4.026834e-04	-6.668434e-03
##	KitchenAbvGr	TotRmsAbvGrd	Fireplaces	GarageCars
##	-7.830963e-03	7.270762e-04	2.321698e-02	2.346675e-02
##	GarageArea	WoodDeckSF	OpenPorchSF	EncPorchSF
##	1.983305e-02	6.241643e-03	8.166466e-03	1.317097e-02
##	PoolArea	MiscVal	age	ageSinceRemodel
##	1.941268e-03	-1.111083e-04	-5.187661e-02	-1.096697e-02
##	MSZoningRH	MSZoningRM	LotShapeIR3	LotShapeReg
##	-3.307184e-03	-2.456559e-02	-2.167316e-03	-1.256580e-03
##	LandContourHLS	LotConfigCulDSac	LandSlopeMod	LandSlopeSev
##	5.603433e-03	5.360943e-03	3.218830e-03	-1.209156e-03
##	NeighborhoodOldTown	NeighborhoodOther	NeighborhoodOther	Condition1Norm
##	-6.718497e-03	-3.094314e-03	2.095832e-03	1.557727e-02

```

##      BldgTypeDuplex      BldgTypeTwnhs      BldgTypeTwnhsE      HouseStyleSLv1
##      -2.030359e-03      -1.061458e-02      -3.343478e-03      9.630373e-04
##      HouseStyleOther      RoofStyleHip      RoofStyleleother      Exterior1stMetalSd
##      -4.150667e-03      7.859146e-04      9.913929e-03      3.892198e-03
##      Exterior1stWd Sdng      Exterior1stOther      Exterior2ndVinylSd      Exterior2ndWd Sdng
##      -2.882589e-03      5.848763e-03      6.424823e-03      4.191821e-03
##      Exterior2ndOther      ExterQualAvg      ExterQualBelowAvg      ExterCondAvg
##      -4.336144e-04      -4.844774e-03      -6.207730e-03      3.657840e-03
##      Foundationother      FoundationPConc      Heatingother      HeatingQCAvg
##      -3.155674e-04      1.777523e-02      3.117311e-03      -6.889703e-03
##      HeatingQCBelowAvg      CentralAirY      KitchenQualAvg      KitchenQualBelowAvg
##      -2.721806e-03      1.042547e-02      -7.274004e-03      -3.839291e-03
##      FunctionalMaj2      FunctionalMin2      FunctionalMod      FunctionalTyp
##      -1.235929e-02      2.701013e-03      -1.480586e-04      1.487099e-02
##      PavedDriveP      PavedDriveY
##      -1.641397e-05      6.425672e-03

```

Model	Notes	Hyperparameters	RMSE	Rsquared
Lasso	caret and elasticnet	Alpha = 1 , Lambda = 0.00167070437878296	0.1010816	0.922609

1 (d) - Model Variants

1 (d, i) - PCR Model

Model Setup

- Uses **numeric** data for Principal Component Analysis
- Then appends the **factor** data to the data *without NULL values*
- Finally, uses `stepAIC()` to best model data
- See interpretation at end

Get cleaned **numeric** and **factor** data frames

Perform PCA

Now we choose number of PC's that explain 75% of the variation

- Note this threshold is just a judgement call. No significance behind 75%

```
## [1] "There are 12 principal components that explain up to 75% of the variation in the data"
```

Join on the factor data

```
df.pcr <- cbind(SalePrice = hd.numericClean$SalePrice, chosenPCs, hd.factorClean)
```

Fit the Model

- Linear model containing:
 - Principal components explaining 75% of variation in **numeric** data
 - Non-null **factor** data
 - *Predicted variable:* `log(SalePrice)`
- Then use `stepAIC()` to identify which variables are actually important for model

```
# Fit data using PC's, non-null factors
fit.pcr <- lm(log(SalePrice) ~ ., data = df.pcr)

# Reduce to only important variables
fit.pcrReduced <- stepAIC(fit.pcr, direction="both")
```

Model	Notes	Hyperparameters	RMSE	Rsquared
PCR	lm	N/A	0.1014132	0.9178239

```
# View results
summary(fit.pcrReduced)
```

```
##
## Call:
## lm(formula = log(SalePrice) ~ PC1 + PC2 + PC3 + PC4 + PC5 + PC6 +
##      PC7 + PC9 + PC12 + MSZoning + LandContour + LotConfig + Condition1 +
##      HouseStyle + RoofStyle + Exterior1st + ExterQual + ExterCond +
##      Foundation + Heating + CentralAir + KitchenQual + Functional +
##      PavedDrive, data = df.pcr)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.69673 -0.05875  0.00159  0.06701  0.31402
```

```

##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  11.7686397  0.0550213  213.892 < 2e-16 ***
## PC1          0.0986702  0.0023721  41.597 < 2e-16 ***
## PC2         -0.0054867  0.0033611  -1.632 0.102923
## PC3         -0.0511019  0.0033405 -15.298 < 2e-16 ***
## PC4         -0.0235630  0.0033654  -7.002 4.78e-12 ***
## PC5         -0.0379678  0.0033373 -11.377 < 2e-16 ***
## PC6         -0.0087772  0.0031369  -2.798 0.005245 **
## PC7          0.0349521  0.0035230   9.921 < 2e-16 ***
## PC9          0.0081982  0.0034207   2.397 0.016739 *
## PC12         0.0181949  0.0036259   5.018 6.23e-07 ***
## MSZoningRH   -0.0659683  0.0398357  -1.656 0.098051 .
## MSZoningRL   -0.0321814  0.0201756  -1.595 0.111031
## MSZoningRM   -0.1125247  0.0218390  -5.152 3.13e-07 ***
## LandContourHLS  0.0744653  0.0266931   2.790 0.005382 **
## LandContourLow  0.0010358  0.0284758   0.036 0.970992
## LandContourLvl -0.0156112  0.0181350  -0.861 0.389548
## LotConfigCulDSac  0.0439742  0.0151677   2.899 0.003827 **
## LotConfigInside  0.0032318  0.0091396   0.354 0.723713
## LotConfigother -0.0071985  0.0192764  -0.373 0.708908
## Condition1Feedr  0.0486306  0.0246766   1.971 0.049047 *
## Condition1Norm   0.0926746  0.0203588   4.552 6.00e-06 ***
## Condition1RR     0.0520491  0.0291863   1.783 0.074851 .
## Condition1Other  0.0244003  0.0311172   0.784 0.433153
## HouseStyle1Story -0.0668074  0.0157227  -4.249 2.36e-05 ***
## HouseStyle2Story -0.0146966  0.0153565  -0.957 0.338797
## HouseStyleSLvl  -0.0270715  0.0203274  -1.332 0.183255
## HouseStyleOther -0.0509964  0.0199468  -2.557 0.010724 *
## RoofStyleHip     0.0152849  0.0093476   1.635 0.102345
## RoofStyleother   0.0970582  0.0246717   3.934 8.96e-05 ***
## Exterior1stMetalSd  0.0290609  0.0125908   2.308 0.021207 *
## Exterior1stVinylSd  0.0256903  0.0113709   2.259 0.024091 *
## Exterior1stWd Sdng -0.0034236  0.0133337  -0.257 0.797416
## Exterior1stOther  0.0314277  0.0115424   2.723 0.006592 **
## ExterQualAvg     -0.0383087  0.0116412  -3.291 0.001036 **
## ExterQualBelowAvg -0.1032930  0.0466318  -2.215 0.026991 *
## ExterCondAvg      0.0222015  0.0116446   1.907 0.056874 .
## ExterCondBelowAvg  0.0060538  0.0322546   0.188 0.851162
## FoundationCBlock  0.0066855  0.0143041   0.467 0.640335
## Foundationother   0.0261987  0.0253505   1.033 0.301651
## FoundationPConc   0.0535857  0.0166839   3.212 0.001363 **
## Heatingother     0.0356275  0.0252882   1.409 0.159204
## CentralAirY      0.0648420  0.0184339   3.518 0.000456 ***
## KitchenQualAvg   -0.0212534  0.0104036  -2.043 0.041339 *
## KitchenQualBelowAvg -0.0412470  0.0257426  -1.602 0.109426
## FunctionalMaj2    -0.2007091  0.0617032  -3.253 0.001183 **
## FunctionalMin1     0.0411613  0.0387696   1.062 0.288646
## FunctionalMin2     0.0371397  0.0377803   0.983 0.325835
## FunctionalMod      0.0145972  0.0446559   0.327 0.743829
## FunctionalTyp      0.1083890  0.0317133   3.418 0.000658 ***
## PavedDriveP      -0.0009913  0.0252404  -0.039 0.968680
## PavedDriveY       0.0503331  0.0161331   3.120 0.001864 **

```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1041 on 949 degrees of freedom
## Multiple R-squared:  0.9219, Adjusted R-squared:  0.9178
## F-statistic: 224.2 on 50 and 949 DF,  p-value: < 2.2e-16
```

View and Interpret Results

Please note all interpretations below are approximate, given the `stepAIC()` uses stochastic modeling.

Model performance evaluation:

- See that around 28 of the variables cannot be explained by random chance, with a probability of 90% or more (see significance codes above)
- Standard errors range from ± 1 -5%, with average around 2%. Larger values may indicate higher uncertainty of the estimated coefficients.
- This model explains around 92% of the variation in the `log(SalePrice)`. See Adjusted R-Squared for reference.
- Note this model may exhibit selection bias, since the data excludes factor data with null values in the variable.
- This model would likely do well for prediction of `log(SalePrice)`, given the small range of standard errors, high adjusted R squared, and number of significant variables. This model would obviously not do well for inference, given we are using principal components that mask the numeric data.

Practical significance evaluation:

- The principal components contribute positively about 20% of the sale price of the home
- Residential Medium Density (`MSZoningRM`) reduces the home price by around 12%, with a standard error of around 2%.
- If the exterior quality is below average (`ExterQualBelowAvg`), it reduces the home price by around 12%, with a standard error of around 5%.
- If the functionality of the home has 2 major deductions (`FunctionalMaj2`), it reduces the home price by around 20%, with a standard error of around 6%. While having typical functionality (`FunctionalTyp`) increases the home sale price by nearly 10%, with a standard error of 3%.
- See other coefficients of the data for other variables.

View Predicted vs. Actuals

Function to compare predicted vs. observed values

```
# Function to compare predicted vs. actual (observed) regression outputs
predictedVsObserved <- function(predicted, observed, modelName, outcomeName = 'Log(SalePrice)') {

  ## Create data set for predicted vs. actuals
  comparison <- data.frame(observed = observed,
                           predicted = predicted) %>%

  # Row index
  mutate(ID = row_number()) %>%

  # Put in single column
  pivot_longer(cols = c('observed', 'predicted'),
               names_to = 'metric',
               values_to = 'value')

  # Plot --- Observed vs. Actuals across all variables in data
  variationScatter <- comparison %>%
    ggplot(aes(x = ID,
               y = value,
               color = metric
               )
           ) +
    geom_point(alpha = 0.5, size = 1) +

    labs(title = 'Variation in Predicted vs. Observed Data',
         subtitle = paste('Model:', modelName),
         x = 'X', y = outcomeName) +
    theme_minimal() + theme(legend.title = element_blank(),
                           legend.position = 'top') +
    scale_color_manual(values = c('grey60', 'palegreen3'))

  print(variationScatter)

  # Limit for x and y axis for scatter of predicted vs. observed
  axisLim = c( min(c(predicted, observed)), max(c(predicted, observed)) )

  # Simple comparison of data
  plot(x = observed,
       y = predicted,
       main = paste(modelName, 'Model - Actual (Observed) vs. Predicted\n'),
       xlab = paste('Observed Values -', outcomeName),
       ylab = paste('Predicted Values -', outcomeName),
       pch = 16,
       cex = 0.75,
       col = alpha('steelblue3', 1/4),
       xlim = axisLim,
       ylim = axisLim
  )
}
```



```

# Add the Predicted vs. actual line
abline(lm(predicted ~ observed), col = 'steelblue3', lwd = 2)
mtext('Predicted ~ Actual', side = 3, adj = 1, col = 'steelblue3')

# Add line for perfectly fit model
abline(0,1, col = alpha('tomato3', 0.8), lwd = 2)
mtext('Perfectly Fit Model', side = 1, adj = 0, col = 'tomato3')
}

```

View results of the PCR Model

- See that the variation in the data is very closely resembled actual by changes in independent variables
- Implication? This model fits its own data well, but what is not know if it can predict out of sample data.
- Note that it the data (blue) deviates slightly from perfect line model (red), indicating that the model is slightly skewed from predicted and actual data.

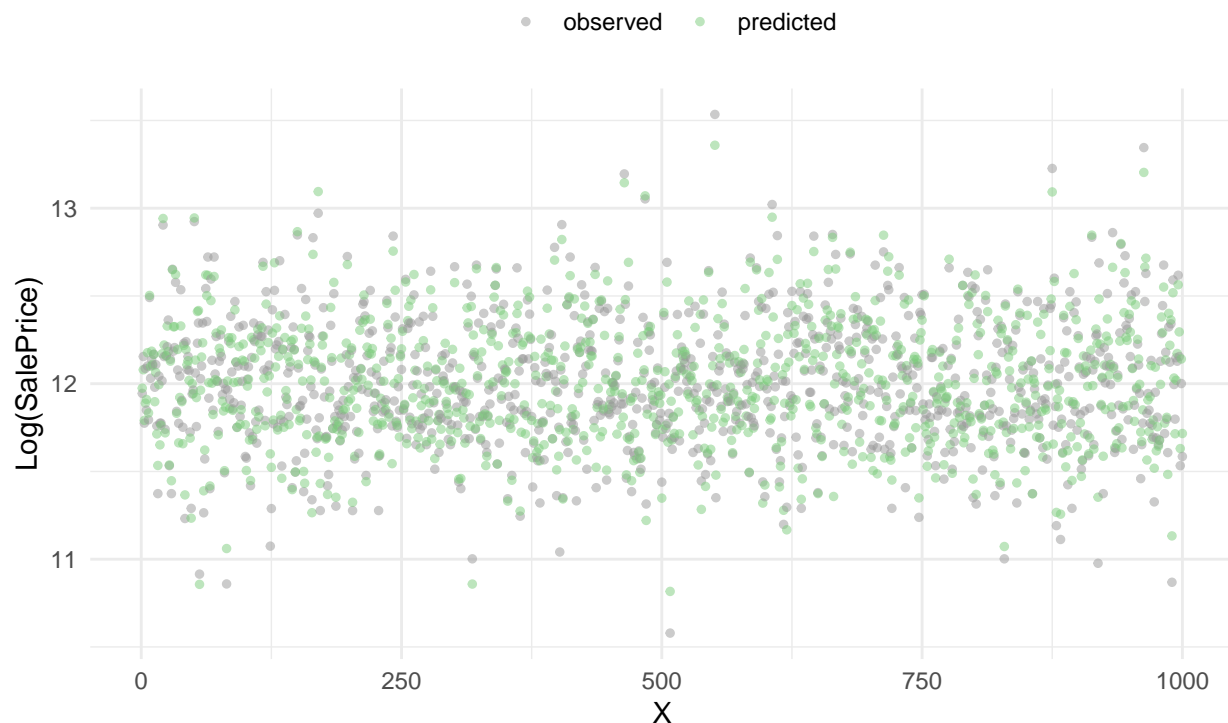
```

# How do the predicted vs. Actuals Compare?
predictedVsObserved(observed = log(df.pcr$SalePrice),
                    predicted = predict(fit.pcrReduced),
                    modelName = 'PCR')

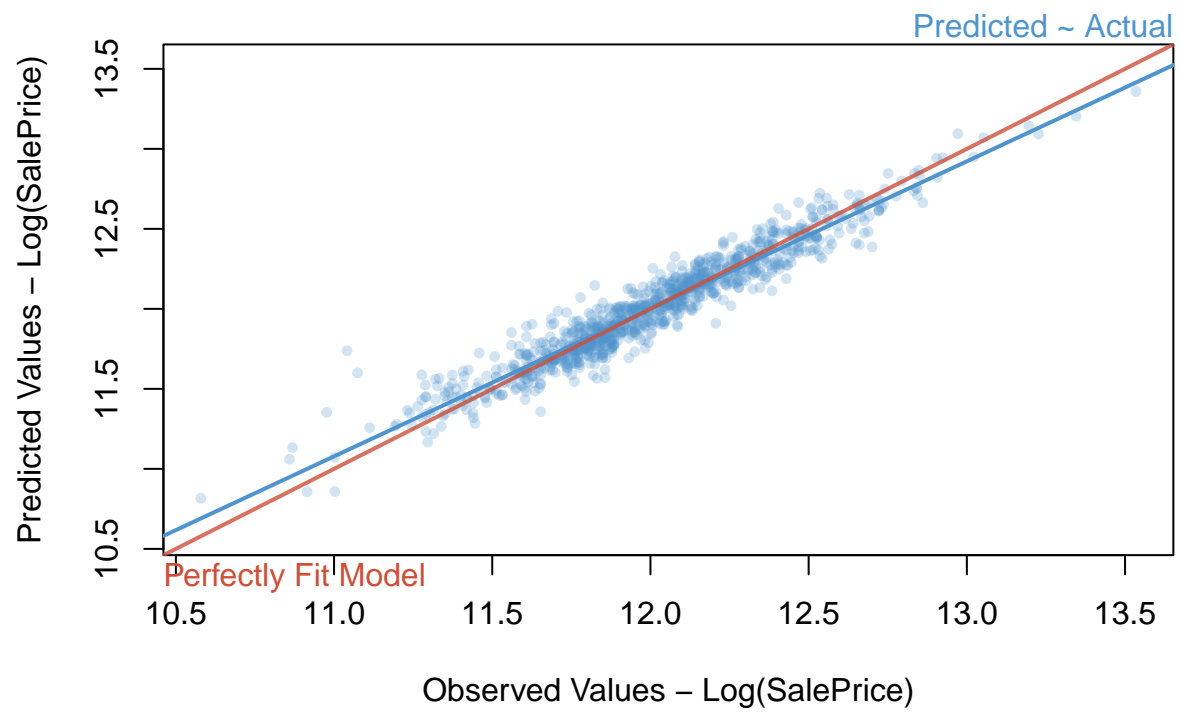
```

Variation in Predicted vs. Observed Data

Model: PCR



PCR Model – Actual (Observed) vs. Predicted



1 (d, ii) - SVR Model

Model Setup

```
ctrl <- trainControl(method = "repeatedcv",
                     number = 5, # 5 fold cross validation
                     repeats = 2 # 2 repeats
                     )

# The data (PMM imputed values, and only whole columns without NA. Does not omit outliers)
df.svm <- cbind(SalePrice = hd.numericClean$SalePrice,
                hd.numericClean, hd.factorClean)
```

Fit the Model

```
# Train and tune the SVM
fit.svm <- train(data = df.svm,
                 log(SalePrice) ~ .,
                 method = "svmRadial",          # Radial kernel
                 tuneLength = 9,                # 9 values of the cost function
                 preProc = c("center","scale"), # Center and scale data
                 trControl = ctrl)
```

View and Interpret Results

- Note all numbers mentioned below are approximate
- See that the R Squared of the model is around 0.86, and RMSE is 0.14
- See that the model predicts the data well.
- Also, note that the model predicts the data with less error than the linear model. See this from the RMSE or scatter plot of predicted values.

Model	Notes	Hyperparameters	RMSE	Rsquared
SVM	caret and svmRadial	C = 4 , Epsilon = 0.1	0.1418915	0.8517242

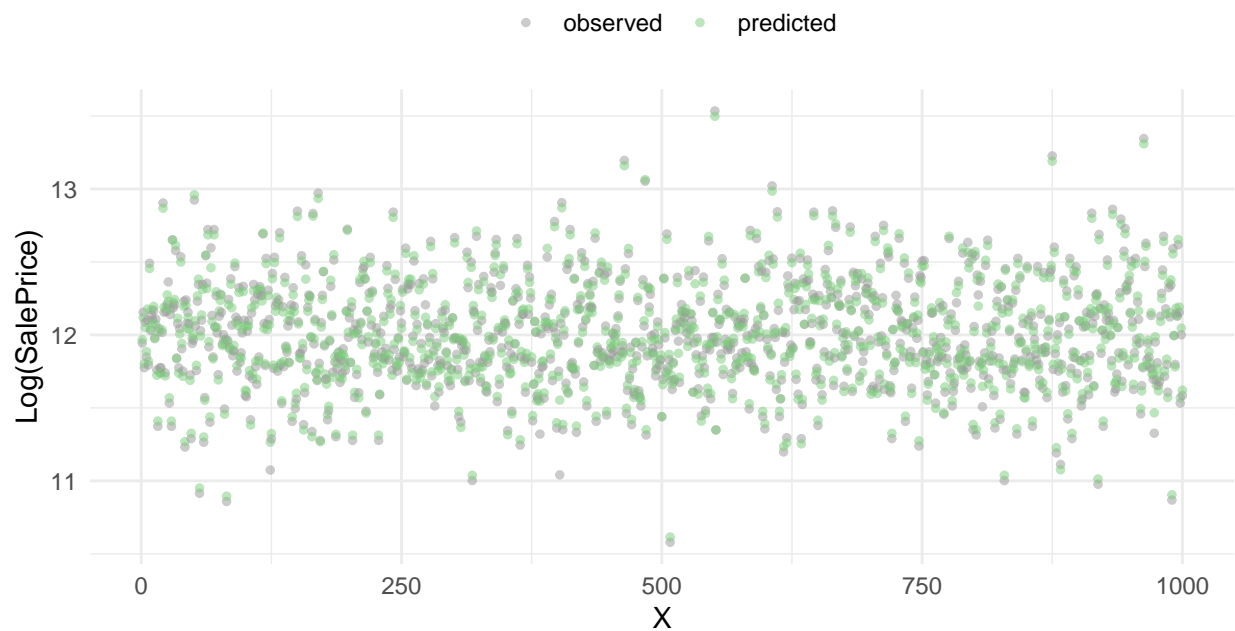
```
# Final model?
fit.svm$finalModel

## Support Vector Machine object of class "ksvm"
##
## SV type: eps-svr (regression)
## parameter : epsilon = 0.1 cost C = 4
##
## Gaussian Radial Basis kernel function.
## Hyperparameter : sigma = 0.00752928895571278
##
## Number of Support Vectors : 671
##
## Objective Function Value : -162.4365
## Training error : 0.012145

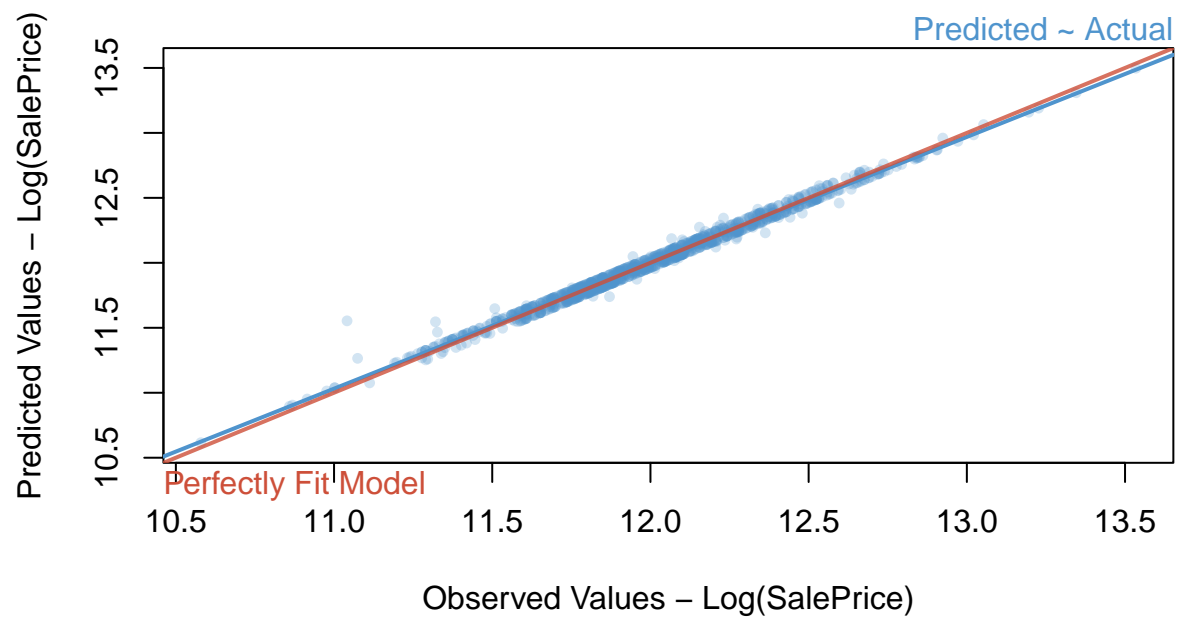
# How do the predicted vs. Actuals Compare?
predictedVsObserved(observed = log(df.svm$SalePrice),
                    predicted = predict(fit.svm, df.svm),
                    modelName = 'SVM')
```

Variation in Predicted vs. Observed Data

Model: SVM



SVM Model – Actual (Observed) vs. Predicted



1 (d, iii) - MARS Model

Fit the Model

```
# Train and tune the MARS model
fit.mars <- train(data = df.svm, # note this is fine since data is the same for this model
  log(SalePrice) ~ .,
  method      = "earth",          # Radial kernel
  tuneLength  = 9,                # 9 values of the cost function
  preProc     = c("center","scale"), # Center and scale data
  trControl   = ctrl
)
```

Model	Notes	Hyperparameters	RMSE	Rsquared
MARS	caret and earth	Degree = 1 , nprune = 17	0.1101934	0.9085685

View and Interpret Results

- See that the model overall performs very well, and in fact performs similarly to the PCR model (in terms of RMSE and Adjusted R Squared).
- Again, unsure if the model would do well for prediction of out of sample data, but fits this data extremely well.

```
# Final model?
fit.mars$finalModel

## Selected 17 of 21 terms, and 10 of 94 predictors (nprune=17)
## Termination condition: RSq changed by less than 0.001 at 21 terms
## Importance: GrLivArea, age, OverallQual, TotalBsmtSF, OverallCond, LotArea, ...
## Number of terms at each degree of interaction: 1 16 (additive model)
## GCV 0.011145    RSS 10.42157    GRSq 0.9155756    RSq 0.9208976

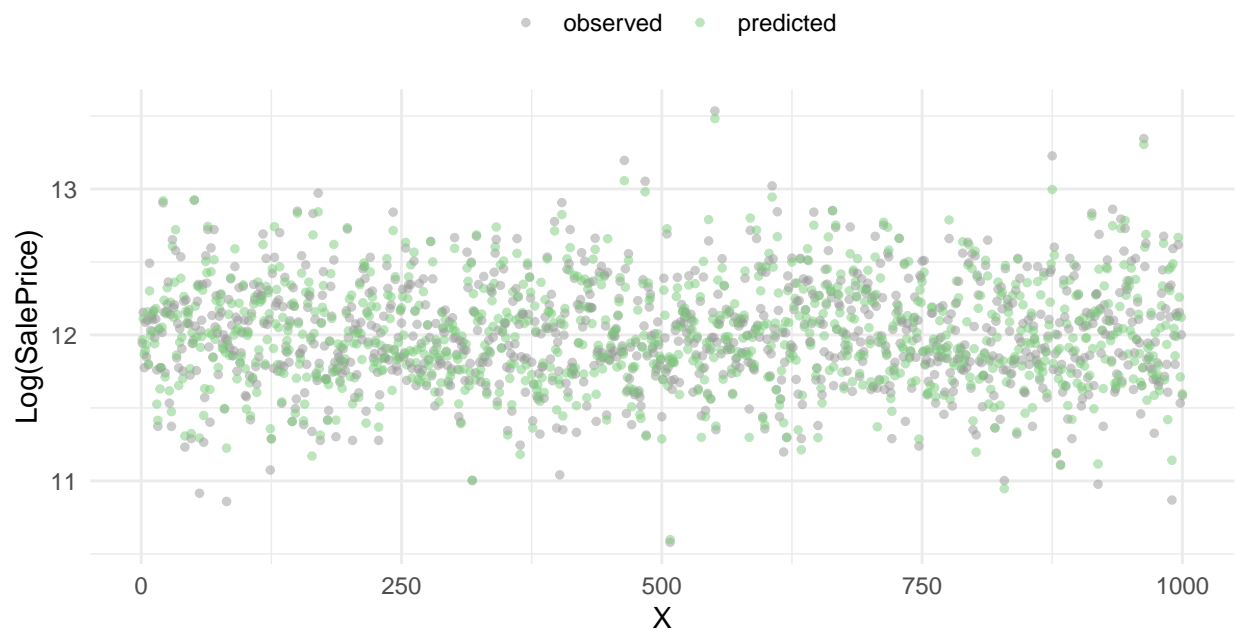
# How do the predicted vs. Actuals Compare?
predicted.mars = fit.mars[["finalModel"]][["fitted.values"]]
colnames(predicted.mars) <- 'predicted'

predictedVsObserved(
  observed = log(df.svm$SalePrice),
  predicted = predicted.mars,
  modelName = 'MARS')

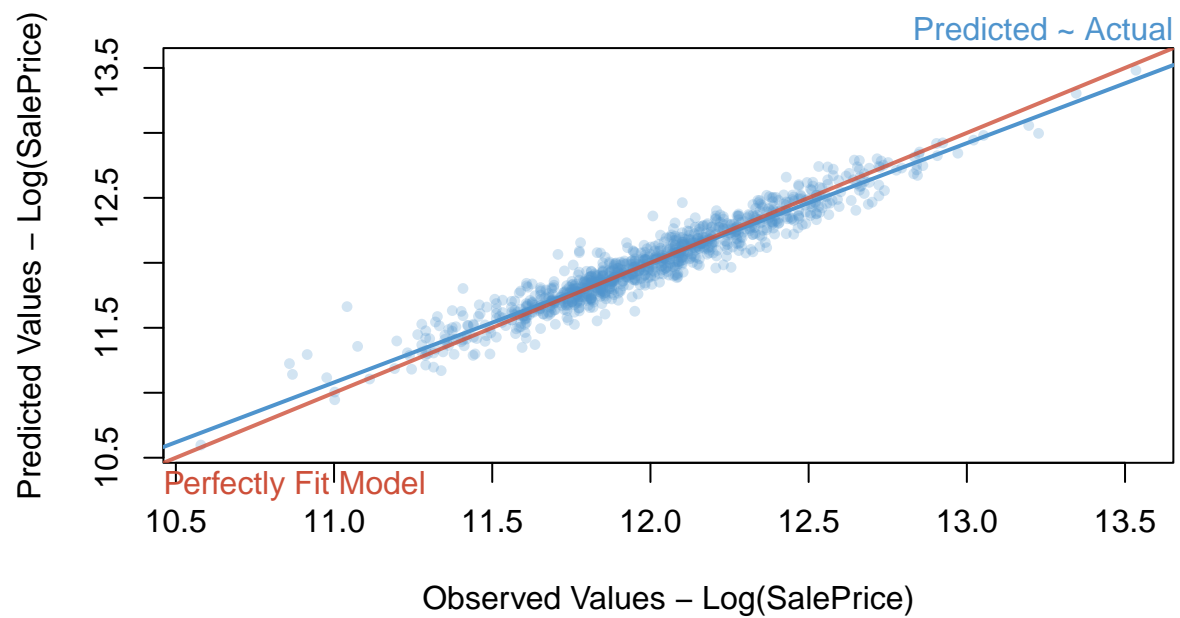
```

Variation in Predicted vs. Observed Data

Model: MARS



MARS Model – Actual (Observed) vs. Predicted



Summary Table of Model Performance

Model	Notes	Hyperparameters	RMSE	Rsquared
OLS	lm	N/A	20948.4222	0.8142
OLS	lm + 2-way interactions	N/A	15929.1310	0.8849
PLS	pls	ncomp = 4	0.1475	0.0218
Lasso	caret and elasticnet	Alpha = 1 , Lambda = 0.00167070437878296	0.1011	0.9226
PCR	lm	N/A	0.1014	0.9178
SVM	caret and svmRadial	C = 4 , Epsilon = 0.1	0.1419	0.8517
MARS	caret and earth	Degree = 1 , nprune = 17	0.1102	0.9086

References

1. https://rpubs.com/staneaurelius/house_price_prediction
2. <https://www.statology.org/partial-least-squares-in-r/>
3. <https://davidalpiaz.github.io/r4sl/elastic-net.html>