

# ISE 5103 Intelligent Data Analytics

## Homework 5 - Modeling

Daniel Carpenter & Sonaxy Mohanty

October 2022

### Contents

Packages . . . . .	2
General Data Prep . . . . .	2
Read Data . . . . .	2
Impute Missing Values with PMM . . . . .	2
Factor Level Collapse - Create Other Bin for Columns over 4 Unique Values . . . . .	5
Remove Outliers from Numeric Data . . . . .	6
<b>1 (a) - OLS Model</b>	<b>7</b>
<b>1 (b) - PLS Model</b>	<b>8</b>
<b>1 (c) - LASSO Model</b>	<b>9</b>
<b>1 (d) - Model Variants</b>	<b>10</b>
1 (d, i) - PCR Model . . . . .	10
Perform PCA analysis to see how Principal components explain variance . . . . .	10
Now, Apply predictions with PCR . . . . .	11
Interpretation of PCR Model . . . . .	13
1 (d, ii) - SVR Model . . . . .	14
1 (d, iii) - MARS Model . . . . .	15

## Packages

```
# Data Wrangling
library(tidyverse)

# Modeling
library(MASS)

# Aesthetics
library(knitr)
library(cowplot) # multiple ggplots on one plot with plot_grid()
library(scales)
library(kableExtra)
```

## General Data Prep

### Read Data

```
hd <- read.csv('housingData.csv') %>%

# creates new variables age, ageSinceRemodel, and ageofGarage, and
dplyr::mutate(age = YrSold - YearBuilt,
              ageSinceRemodel = YrSold - YearRemodAdd,
              ageofGarage = ifelse(is.na(GarageYrBlt), age, YrSold - GarageYrBlt)) %>%

# removes the columns used in above the calculations
dplyr::select(!c(Id, MSSubClass, MiscVal, YrSold,
                 MoSold, YearBuilt, YearRemodAdd))

# Convert all character data to factor
hd[sapply(hd, is.character)] <-
  lapply(hd[sapply(hd, is.character)], as.factor)
```

### Impute Missing Values with PMM

Make dataset of **numeric** variables

```
hd.numericRaw <- hd %>%

#selecting all the numeric data
dplyr::select_if(is.numeric) %>%

#converting the dataframe to tibble
as_tibble()
```

Make dataset of **factor** variables

```

hd.factorRaw <- hd %>%

#selecting all the numeric data
dplyr::select_if(is.numeric) %>%

#converting the dataframe to tibble
as_tibble()

```

For each column with missing data, impute missing values with PMM

- Done with function `imputeWithPMM()` function
- Applies function via `dplyr` logic
- Note `seeImputation()` function to visualize the imputation from prior homework 4, not shown for simplicity in viewing

Create function to impute via PMM

```

imputeWithPMM <- function(colWithMissingData) {

# Using the mice package
suppressMessages(library(mice))

# Discover the missing rows
isMissing <- is.na(colWithMissingData)

# Create data frame to pass to PMM imputation function from mic package
df <- data.frame(x      = rexp(length(colWithMissingData)), # meaningless x to help show variation
                 y      = colWithMissingData,
                 missing = isMissing)

# imputation by PMM
df[isMissing, "y"] <- mice.impute.pmm( df$y,
                                       !df$missing,
                                       df$x)

return(df$y)
}

```

Apply PMM function to numeric data containing null values

```

# Data to store imputed values with PMM method
hd.Imputed <- hd

# Which columns has NA's?
colNamesWithNulls <- colnames(hd.numericRaw[ , colSums(is.na(hd.numericRaw)) != 0])
colNamesWithNulls

```

```
## [1] "LotFrontage" "MasVnrArea" "GarageYrBlt"
```

```

numberOfColsWithNulls = length(colNamesWithNulls)

# For each of the numeric columns with null values
for (colWithNullsNum in 1:numberOfColsWithNulls) {

  # The name of the column with null values
  nameOfThisColumn <- colNamesWithNulls[colWithNullsNum]

  # Get the actual data of the column with nulls
  colWithNulls <- hd[, nameOfThisColumn]

  # Impute the missing values with PMM
  imputedValues <- imputeWithPMM(colWithNulls)

  # Now store the data in the original new frame
  hd.Imputed[, nameOfThisColumn] <- imputedValues

  # Save a visualization of the imputation
  pmmVisual <- seeImputation(data.frame(y = colWithNulls),
                             data.frame(y = imputedValues),
                             nameOfThisColumn )

  fileToSave = paste0('OutputPMM/Imputation_With_PMM_', nameOfThisColumn, '.pdf')
  print(paste0('For imputation results of ', nameOfThisColumn, ', see ', fileToSave))
  ggsave(pmmVisual, filename = fileToSave,
         height = 11, width = 8.5)
}

```

```
## [1] "For imputation results of LotFrontage, see OutputPMM/Imputation_With_PMM_LotFrontage.pdf"
```

```
## [1] "For imputation results of MasVnrArea, see OutputPMM/Imputation_With_PMM_MasVnrArea.pdf"
```

```
## [1] "For imputation results of GarageYrBlt, see OutputPMM/Imputation_With_PMM_GarageYrBlt.pdf"
```

## Factor Level Collapse - Create Other Bin for Columns over 4 Unique Values

```
hd.Cleaned <- hd.Imputed # For final cleaned data

# Get list of factors and the number of unique values
factorCols <- as.data.frame(t(hd.factorRaw %>% summarise_all(n_distinct)))

# We are going to factor collapse factor columns with more than 4 columns
# So there will be 4 of the original, and 1 containing 'other'
# This is the threshold
factorThreshold = 4

# Get a list of the factors we are going to collapse
colsWithManyFactors <- rownames(factorCols %>% filter(V1 > factorThreshold))

# Show a summary of how many factors will be collapsed
numberOfColsWithManyFactors = length(colsWithManyFactors)
paste('Before cleaning, there are', numberOfColsWithManyFactors, 'factor columns with more than',
      factorThreshold, 'unique values')
```

```
## [1] "Before cleaning, there are 14 factor columns with more than 4 unique values"
```

```
# Collapse the affected factors in the original data (the one that already has imputation)

## for each factor column that we are about to collapse
for (collapsedColNum in 1:numberOfColsWithManyFactors) {

  # The name of the column with null values
  nameOfThisColumn <- colsWithManyFactors[collapsedColNum]

  # Get the actual data of the column with nulls
  colWithManyFactors <- hd[, nameOfThisColumn]

  # lumps all levels except for the n most frequent
  hd.Cleaned[, nameOfThisColumn] <- fct_lump_n(colWithManyFactors,
                                              n=factorThreshold)
}

# Check to see if the factor lumping worked
factorColsCleaned <- t(hd.Cleaned %>%
  select_if(is.factor) %>%
  summarise_all(n_distinct))
paste('After cleaning, there are', sum(factorColsCleaned > factorThreshold, na.rm = TRUE),
      "columns with more than", factorThreshold, "unique values (omitting NA's)")
```

```
## [1] "After cleaning, there are 14 columns with more than 4 unique values (omitting NA's)"
```

## Remove Outliers from Numeric Data

- Since there are so many outliers, we are only going to remove some outliers
- If you count the number of outliers by column, the 75% of columns contain less than 50 outliers.
- However, some contain up to 200. Since remove ALL outliers would reduce the size of the data to less than 300 observations, we are removing up to 50 per column.

```
hd.CleanedNoOutliers <- hd.Cleaned

# Remove up to 75% of the outliers in the dataset
# this is the 3rd quartile of number of outliers.
k_outliers = 50
numOutliers = c() # to store the number of outliers per column

theColNames <- colnames(hd.Cleaned)

for (colNum in 1:ncol(hd.Cleaned)) {

  theCol <- hd.Cleaned[, colNum]
  nrowBefore = length(theCol)
  colName <- theColNames[colNum]

  # Only consider numeric
  if (is.numeric(theCol)) {

    # Identify the outliers in the column
    # Source: https://www.geeksforgeeks.org/remove-outliers-from-data-set-in-r/
    columnOutliers <- boxplot.stats(hd.CleanedNoOutliers[, colNum])$out
    numOutliers <- c(numOutliers, length(columnOutliers))

    # Now remove k outliers from the column
    if (length(columnOutliers) < k_outliers) {

      hd.CleanedNoOutliers <- hd.CleanedNoOutliers %>%

        # If this syntax looks weird, it is just referencing a column in the
        # dataset using dplyr piping. See below for more info:
        # https://stackoverflow.com/questions/48062213/dplyr-using-column-names-as-function-arguments
        # https://stackoverflow.com/questions/72673381/column-names-as-variables-in-dplyr-select-v-filter
        filter( !( get({colName}) ) %in% columnOutliers ) )
    }
  }
}

paste0('Of the columns with outliers, removed up to 75th percentile of num. outliers.')

## [1] "Of the columns with outliers, removed up to 75th percentile of num. outliers."

paste0('See that the 75th percentile of columns with outliers contain ',
       paste0(summary(numOutliers)[5]), ' outliers')

## [1] "See that the 75th percentile of columns with outliers contain 50.25 outliers"
```

## 1 (a) - OLS Model

## 1 (b) - PLS Model



## 1 (c) - LASSO Model

## 1 (d) - Model Variants

### 1 (d, i) - PCR Model

Perform PCA analysis to see how Principal components explain variance

- Uses `numeric` data for Principal Component Analysis
- Then appends the `factor` data to the data *without NULL values*
- Finally, uses `stepAIC()` to best model data
- See interpretation at end

Get cleaned `numeric` and `factor` data frames

```
# After cleaning, two datasets that contain..

## Numeric data -----
hd.numericClean <- hd.Cleaned %>% select_if(is.numeric)

## Factors -----
hd.factorClean <- hd.Cleaned %>% dplyr::select(where(is.factor))

# Removing any columns with NA
removeColsWithNA <- function(df) {
  return( df[ , colSums(is.na(df)) == 0] )
}
hd.factorClean <- removeColsWithNA(hd.factorClean)

paste('Num. factor cols. removed due to null values:',
      ncol(hd.Cleaned %>% dplyr::select(where(is.factor))) - ncol(hd.factorClean) )
```

```
## [1] "Num. factor cols. removed due to null values: 16"
```

```
paste(ncol(hd.factorClean), 'factor cols. remain')
```

```
## [1] "22 factor cols. remain"
```

Perform PCA

```
# Principal component analysis on numeric data
pc.house <- prcomp(hd.numericClean %>% dplyr::select(-SalePrice), # do not include response var
                  center = TRUE, # Mean centered
                  scale = TRUE # Z-Score standardized
                  )

# See first 10 cumulative proportions
pc.house.summary <- summary(pc.house)
pc.house.summary$importance[, 1:10]
```

```
##               PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  2.657147 1.84832 1.609823 1.393088 1.171847 1.107559
## Proportion of Variance 0.227760 0.11020 0.083600 0.062600 0.044300 0.039570
## Cumulative Proportion 0.227760 0.33796 0.421560 0.484160 0.528460 0.568030
##               PC7      PC8      PC9      PC10
## Standard deviation  1.060806 1.034741 1.008248 1.005577
## Proportion of Variance 0.036300 0.034540 0.032790 0.032620
## Cumulative Proportion 0.604330 0.638870 0.671660 0.704280
```

Now we choose number of PC's that explain 75% of the variation

- Note this threshold is just a judgement call. No significance behind 75%

```
cumPropThreshold = 0.75 # The threshold

numPCs <- sum(pc.house.summary$importance['Cumulative Proportion', ] < cumPropThreshold)
paste0('There are ', numPCs, ' principal components that explain up to ', cumPropThreshold*100,
      '% of the variation in the data')
```

```
## [1] "There are 11 principal components that explain up to 75% of the variation in the data"
```

```
chosenPCs <- as.data.frame(pc.house$x[, 1:numPCs])
```

Join on the factor data

```
df.pcr <- cbind(SalePrice = hd.numericClean$SalePrice, chosenPCs, hd.factorClean)
```

## Now, Apply predictions with PCR

- Linear model containing:
  - Principal components explaining 75% of variation in numeric data
  - Non-null factor data
  - *Predicted variable:* log(SalePrice)
- Then use `stepAIC()` to identify which variables are actually important for model

```
# Fit data using PC's, non-null factors
fit.pcr <- lm(log(SalePrice) ~ ., data = df.pcr)

# Reduce to only important variables
fit.pcrReduced <- stepAIC(fit.pcr, direction="both")
```

```
# View results
summary(fit.pcrReduced)
```

```
##
## Call:
## lm(formula = log(SalePrice) ~ PC1 + PC2 + PC3 + PC4 + PC5 + PC7 +
##      PC8 + PC9 + MSZoning + LandContour + LotConfig + Condition1 +
```

```

##      BldgType + HouseStyle + RoofStyle + Exterior1st + ExterQual +
##      Foundation + CentralAir + KitchenQual + Functional + PavedDrive,
##      data = df.pcr)
##
## Residuals:
##      Min        1Q      Median        3Q        Max
## -0.67073 -0.06094  0.00279  0.06714  0.31587
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    11.816119   0.053210  222.067 < 2e-16 ***
## PC1              0.098443   0.002399   41.033 < 2e-16 ***
## PC2              0.004896   0.003404    1.438 0.150693
## PC3             -0.053829   0.003397  -15.845 < 2e-16 ***
## PC4             -0.018459   0.003634   -5.079 4.57e-07 ***
## PC5              0.053802   0.003954   13.606 < 2e-16 ***
## PC7              0.009082   0.003407    2.665 0.007819 **
## PC8             -0.013184   0.003562   -3.701 0.000227 ***
## PC9              0.007889   0.003647    2.163 0.030770 *
## MSZoningRH      -0.055675   0.040153   -1.387 0.165897
## MSZoningRL      -0.036339   0.020456   -1.776 0.075990 .
## MSZoningRM      -0.113188   0.022046   -5.134 3.44e-07 ***
## LandContourHLS    0.082312   0.026970    3.052 0.002336 **
## LandContourLow   -0.004804   0.028840   -0.167 0.867737
## LandContourLvl   -0.005914   0.018233   -0.324 0.745736
## LotConfigCulDSac  0.044585   0.015183    2.937 0.003399 **
## LotConfigInside  0.006491   0.009151    0.709 0.478304
## LotConfigother   -0.002854   0.019360   -0.147 0.882818
## Condition1Feedr   0.052676   0.024908    2.115 0.034706 *
## Condition1Norm    0.093188   0.020614    4.521 6.95e-06 ***
## Condition1RR      0.053171   0.029508    1.802 0.071874 .
## Condition1Other   0.019094   0.031505    0.606 0.544619
## BldgType2fmCon    0.035597   0.027343    1.302 0.193278
## BldgTypeDuplex    0.056117   0.026551    2.114 0.034814 *
## BldgTypeTwnhs    -0.046161   0.022101   -2.089 0.037009 *
## BldgTypeTwnhsE   -0.003571   0.015351   -0.233 0.816109
## HouseStyle1Story -0.066389   0.015610   -4.253 2.32e-05 ***
## HouseStyle2Story -0.010274   0.015247   -0.674 0.500584
## HouseStyleSLvl   -0.033754   0.020576   -1.640 0.101247
## HouseStyleOther  -0.055199   0.020248   -2.726 0.006526 **
## RoofStyleHip      0.017808   0.009405    1.893 0.058599 .
## RoofStyleother    0.105921   0.024857    4.261 2.24e-05 ***
## Exterior1stMetalSd 0.024019   0.012770    1.881 0.060287 .
## Exterior1stVinylSd 0.023388   0.011479    2.037 0.041892 *
## Exterior1stWd Sdng -0.005424   0.013439   -0.404 0.686589
## Exterior1stOther  0.033857   0.011735    2.885 0.004000 **
## ExterQualAvg     -0.037753   0.011765   -3.209 0.001377 **
## ExterQualBelowAvg -0.135204   0.045739   -2.956 0.003194 **
## FoundationCBlock  0.002849   0.014310    0.199 0.842210
## Foundationother   0.025238   0.025736    0.981 0.327007
## FoundationPConc   0.053678   0.016765    3.202 0.001411 **
## CentralAirY       0.055811   0.017309    3.224 0.001306 **
## KitchenQualAvg   -0.025722   0.010425   -2.467 0.013789 *
## KitchenQualBelowAvg -0.048222   0.025028   -1.927 0.054317 .

```

```
## FunctionalMaj2      -0.223632    0.062105   -3.601 0.000334 ***
## FunctionalMin1       0.024115    0.039218    0.615 0.538782
## FunctionalMin2       0.022843    0.038109    0.599 0.549033
## FunctionalMod        -0.006688    0.044736   -0.149 0.881193
## FunctionalTyp        0.088687    0.032056    2.767 0.005774 **
## PavedDriveP         -0.010921    0.025435   -0.429 0.667746
## PavedDriveY          0.045561    0.016324    2.791 0.005360 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1049 on 949 degrees of freedom
## Multiple R-squared:  0.9207, Adjusted R-squared:  0.9165
## F-statistic: 220.3 on 50 and 949 DF,  p-value: < 2.2e-16
```

## Interpretation of PCR Model

*Please note all interpretations below are approximate, given the `stepAIC()` uses stochastic modeling.*

### Model performance evaluation:

- See that around 28 of the variables cannot be explained by random chance, with a probability of 90% or more (see significance codes above)
- Standard errors range from  $\pm 1$ -5%, with average around 2%. Larger values may indicate higher uncertainty of the estimated coefficients.
- This model explains around 92% of the variation in the `log(SalePrice)`. See Adjusted R-Squared for reference.
- Note this model may exhibit selection bias, since the data excludes factor data with null values in the variable.
- This model would likely do well for prediction of `log(SalePrice)`, given the small range of standard errors, high adjusted R squared, and number of significant variables. This model would obviously not do well for inference, given we are using principal components that mask the numeric data.

### Practical significance evaluation:

- The principal components contribute positively about 20% of the sale price of the home
- Residential Medium Density (`MSZoningRM`) reduces the home price by around 12%, with a standard error of around 2%.
- If the exterior quality is below average (`ExterQualBelowAvg`), it reduces the home price by around 12%, with a standard error of around 5%.
- If the functionality of the home has 2 major deductions (`FunctionalMaj2`), it reduces the home price by around 20%, with a standard error of around 6%. While having typical functionality (`FunctionalTyp`) increases the home sale price by nearly 10%, with a standard error of 3%.
- See other coefficients of the data for other variables.

## 1 (d, ii) - SVR Model

1 (d, iii) - MARS Model